



learnbyexample

COMMAND LINE TEXT PROCESSING with GNU COREUTILS

an example based guide

**Sundeep
Agarwal**

Table of contents

Preface	6
Prerequisites	6
Conventions	6
Acknowledgements	6
Feedback and Errata	6
Author info	7
License	7
Book version	7
Introduction	8
Installation	8
Documentation	8
cat and tac	9
Creating text files	9
Concatenate files	10
Accepting stdin data	10
Squeeze consecutive empty lines	11
Prefix line numbers	11
Viewing special characters	12
Useless use of cat	13
tac	14
Customize line separator for tac	15
head and tail	17
Leading and trailing lines	17
Excluding the last N lines	18
Starting from Nth line	18
Multiple input files	19
Byte selection	19
Range of lines	20
NUL separator	20
Further Reading	20
tr	21
Translation	21
Different length sets	21
Escape sequences and character sets	22
Deleting characters	23
Complement	23
Squeeze	23
cut	25
Individual field selections	25
Field ranges	25
Input field delimiter	26
Output field delimiter	26
Complement	26
Suppress lines without delimiters	27
Character selections	27

NUL separator	28
Alternatives	28
seq	29
Integer sequences	29
Floating-point sequences	30
Customizing separator	30
Leading zeros	31
printf style formatting	31
Limitations	31
shuf	33
Randomize input lines	33
Limit output lines	33
Repeated lines	34
Specify input as arguments	34
Generate random numbers	35
Specifying output file	36
NUL separator	36
paste	37
Concatenating files column wise	37
Interleaving lines	38
Multiple columns from single input	38
Multicharacter delimiters	39
Serialize	40
NUL separator	40
pr	41
Columnate	41
Customizing page width	42
Concatenating files column wise	43
Miscellaneous	44
fold and fmt	45
fold	45
fmt	46
sort	48
Default sort and Collating order	48
Ignoring headers	49
Dictionary sort	49
Reversed order	50
Numeric sort	50
Human numeric sort	51
Version sort	52
Random sort	53
Unique sort	53
Column sort	54
Character positions within columns	56
Debugging	57
Check if sorted	58

Specifying output file	58
Merge sort	59
NUL separator	59
Further Reading	60
uniq	61
Retain single copy of duplicates	61
Duplicates only	61
Unique only	62
Grouping similar lines	62
Prefix count	63
Ignoring case	64
Partial match	64
Specifying output file	65
NUL separator	65
Alternatives	65
comm	66
Three column output	66
Suppressing columns	67
Duplicate lines	68
NUL separator	68
Alternatives	68
join	69
Default join	69
Non-matching lines	70
Change field separator	70
Files with headers	71
Change key field	71
Customize output field list	72
Same number of output fields	72
Set operations	73
NUL separator	75
Alternatives	75
nl	76
Default numbering	76
Number formatting	76
Customize width	77
Customize separator	77
Starting number and increment	77
Section wise numbering	78
Section numbering criteria	81
wc	82
Line, word and byte counts	82
Individual counts	82
Multiple files	83
Character count	83
Longest line length	83
Corner cases	84

split	86
Default split	86
Change number of lines	86
Split by byte count	87
Divide based on file size	88
Interleaved lines	90
Custom line separator	90
Customize filenames	91
Exclude empty files	92
Process parts through another command	92
csplit	94
Split on Nth line	94
Split on regexp	95
Regexp offset	96
Repeat split	97
Keep files on error	98
Suppress matched lines	99
Exclude empty files	100
Customize filenames	101
expand and unexpand	103
Default expand	103
Expand only initial tabs	103
Customize tab stop width	104
Default unexpand	105
Unexpand all blanks	106
Change tab stop width	106
basename and dirname	108
Extract filename from path	108
Remove file extension	108
Remove filename from path	109
Multiple arguments	109
Combining basename and dirname	109
NUL separator	109
What next?	111

Preface

You might be already aware of popular coreutils commands like `head` , `tail` , `tr` , `sort` , etc. This book will teach you more than twenty of such specialized text processing tools provided by the `GNU coreutils` package.

My [Command Line Text Processing](#) repo includes chapters on some of these coreutils commands. Those chapters have been significantly edited for this book and new chapters have been added to cover more commands.

Prerequisites

Prior experience working with command line and `bash` shell, should know concepts like file redirection, command pipeline and so on.

If you are new to the world of command line, check out my curated resources on [Linux CLI and Shell scripting](#) before starting this book.

Conventions

- The examples presented here have been tested on `GNU bash` shell and **version 8.30** of the `GNU coreutils` package.
- Code snippets shown are copy pasted from `bash` shell and modified for presentation purposes. Some commands are preceded by comments to provide context and explanations. Blank lines have been added to improve readability, only `real` time is shown for speed comparisons and so on.
- Unless otherwise noted, all examples and explanations are meant for **ASCII** characters.
- External links are provided throughout the book for you to explore certain topics in more depth.
- The [cli_text_processing_coreutils](#) repo has all the code snippets, example files and other details related to the book. If you are not familiar with `git` command, click the **Code** button on the webpage to get the files.

Acknowledgements

- [/r/commandline/](#), [/r/linux4noobs/](#) and [/r/linux/](#) — helpful forums
- [stackoverflow](#) and [unix.stackexchange](#) — for getting answers on pertinent questions related to cli tools
- [tex.stackexchange](#) — for help on `pandoc` and `tex` related questions
- [canva](#) — cover image
- [Warning](#) and [Info](#) icons by [Amada44](#) under public domain
- [pngquant](#) and [svgcleaner](#) for optimizing images

Feedback and Errata

I would highly appreciate if you'd let me know how you felt about this book, it would help to improve this book as well as my future attempts. Also, please do let me know if you spot any error or typo.

Issue Manager: https://github.com/learnbyexample/cli_text_processing_coreutils/issues

E-mail: learnbyexample.net@gmail.com

Twitter: https://twitter.com/learn_byexample

Author info

Sundeeep Agarwal is a freelance trainer, author and mentor. His previous experience includes working as a Design Engineer at Analog Devices for more than 5 years. You can find his other works, primarily focused on Linux command line, text processing, scripting languages and curated lists, at <https://github.com/learnbyexample>. He has also been a technical reviewer for [Command Line Fundamentals](#) book and video course published by Packt.

List of books: <https://learnbyexample.github.io/books/>

License

This work is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License](#)

Code snippets are available under [MIT License](#)

Resources mentioned in Acknowledgements section above are available under original licenses.

Book version

1.0

See [Version_changes.md](#) to track changes across book versions.

Introduction

I've been using Linux since 2007, but it took me ten more years to *really* explore coreutils for my [Command Line Text Processing](#) repository.

Any beginner learning Linux command line tools would come across `cat` within the first week. Sooner or later, they'll come to know popular text processing tools like `grep` , `head` , `tail` , `tr` , `sort` , etc. If you were like me, you'd come across `sed` and `awk` , shudder at their complexity and prefer to use a scripting language like Perl and text editors like Vim instead (don't worry, I've already corrected that mistake).

Knowing power tools like `grep` , `sed` and `awk` can help solve most of your text processing needs. So, why would you want to learn text processing tools from the coreutils package? The biggest motivation would be faster execution since these tools are optimized for the use cases they solve. And there's always the advantage of not having to write code (and test that solution) if there's an existing tool to solve the problem.

This book will teach you more than twenty of such specialized text processing tools provided by the `GNU coreutils` package. Plenty of examples are provided to make it easier to understand a particular tool and its various features.

Writing a book always has a few pleasant surprises for me. For this one, it was discovering a `sort` option for calendar months, regular expression based features of `tac` and `nl` commands, etc.

Installation

On a GNU/Linux based OS, you are most likely to already have GNU coreutils installed. This book covers **version 8.30** of the coreutils package. To install a newer/particular version, see [Coreutils download](#) section for links and details.

If you are not using a Linux distribution, you may be able to access coreutils using these options:

- [WSL](#)
- [brew](#)

Documentation

It is always a good idea to know where to find the documentation. From the command line, you can use the `man` and `info` commands for brief manual and full documentation respectively. I prefer using the [online GNU coreutils manual](#) which feels much easier to use and navigate.

See also:

- [Release notes](#) — overview of changes and bug fixes between versions
- [Bug list](#)
- [Common options](#)
- [FAQ](#)
- [Gotchas](#)

cat and tac

`cat` derives its name from **concatenation** and provides other nifty options too.

`tac` helps you to reverse the input line wise, usually used for further text processing.

Creating text files

Yeah, `cat` can be used to write contents to a file by typing them from the terminal itself. If you invoke `cat` without providing file arguments or `stdin` from a pipe, it will wait for you to type the content. After you are done typing all the text you want to save, press Enter key and then `Ctrl+d` key combination. If you don't want the last line to have a newline character, press `Ctrl+d` twice instead of Enter and `Ctrl+d`. See also [unix.stackexchange: difference between Ctrl+c and Ctrl+d](#).

```
# press Enter key and Ctrl+d after typing all the required characters
$ cat > greeting.txt
Hi there
Have a nice day
```

In the above example, the output of `cat` is redirected to a file named `greeting.txt`. If you don't redirect the `stdout`, each line will be echoed as you type. You can check the contents of the file you just created by using `cat` again.

```
$ cat greeting.txt
Hi there
Have a nice day
```

Here Documents is another popular way to create such files. Especially in shell scripts, since pressing `Ctrl+d` interactively won't be possible. Here's an example:

```
# > and a space at the start of lines are only present in interactive mode
# don't type them in a shell script
# EOF is typically used as the identifier
$ cat << 'EOF' > fruits.txt
> banana
> papaya
> mango
> EOF

$ cat fruits.txt
banana
papaya
mango
```

The termination string is enclosed in single quotes to prevent parameter expansion, command substitution, etc. You can also use `\string` for this purpose. If you use `<<-` instead of `<<`, you can use leading tab characters for indentation purposes. See [bash manual: Here Documents](#) and [stackoverflow: here-documents](#) for more details.



Note that creating files as shown above isn't restricted to `cat`, it can be applied to any command waiting for `stdin`.

```
# 'tr' converts lowercase alphabets to uppercase in this example
$ tr 'a-z' 'A-Z' << 'end' > op.txt
> hi there
> have a nice day
> end

$ cat op.txt
HI THERE
HAVE A NICE DAY
```

Concatenate files

Here's some examples to showcase `cat`'s main utility. One or more files can be given as arguments.



Visit the [cli_text_processing_coreutils](#) repo to get all the example files used in this book.

```
$ cat greeting.txt fruits.txt nums.txt
Hi there
Have a nice day
banana
papaya
mango
3.14
42
1000
```

To save the output of concatenation, use your shell's redirection features.

```
$ cat greeting.txt fruits.txt nums.txt > op.txt

$ cat op.txt
Hi there
Have a nice day
banana
papaya
mango
3.14
42
1000
```

Accepting stdin data

You can represent `stdin` data using `-` as a file argument. If file arguments are not present, `cat` will read from `stdin` data if present or wait for interactive input as seen earlier.

```
# only stdin (- is optional in this case)
$ echo 'apple banana cherry' | cat
apple banana cherry
```

```
# both stdin and file arguments
$ echo 'apple banana cherry' | cat greeting.txt -
Hi there
Have a nice day
apple banana cherry

# here's an example without newline character at the end of first input
$ printf 'Some\nNumbers' | cat - nums.txt
Some
Numbers3.14
42
1000
```

Squeeze consecutive empty lines

As mentioned before, `cat` provides many features beyond concatenation. Consider this sample `stdin` data:

```
$ printf 'hello\n\n\nworld\n\nhave a nice day\n'
hello

world

have a nice day
```

You can use the `-s` option to squeeze consecutive empty lines to a single empty line. If present, leading and trailing empty lines will also be squeezed, won't be completely removed. You can modify the below example to test it out.

```
$ printf 'hello\n\n\nworld\n\nhave a nice day\n' | cat -s
hello

world

have a nice day
```

Prefix line numbers

The `-n` option will prefix line number and a tab character to each input line. The line numbers are right justified to occupy a minimum of `6` characters, with space as the filler.

```
$ cat -n greeting.txt fruits.txt nums.txt
  1 Hi there
  2 Have a nice day
  3 banana
  4 papaya
  5 mango
  6 3.14
  7 42
  8 1000
```

Use `-b` option instead of `-n` option if you don't want empty lines to be numbered.

```
# -n option numbers all the input lines
$ printf 'apple\n\nbanana\n\ncherry\n' | cat -n
 1 apple
 2
 3 banana
 4
 5 cherry

# -b option numbers only the non-empty input lines
$ printf 'apple\n\nbanana\n\ncherry\n' | cat -b
 1 apple

 2 banana

 3 cherry
```



Use `nl` command if you want more customization options for numbering.

Viewing special characters

Characters like backspace and carriage return will mangle the contents if viewed naively on the terminal. Characters like NUL won't even be visible. You can use the `-v` option to show such characters using the caret notation (see [wikipedia: Control code chart](#) for details). See [this unix.stackexchange thread](#) for non-ASCII character examples.

```
# example for backspace and carriage return
$ printf 'car\bt\nbike\rp\n'
cat
pike
$ printf 'car\bt\nbike\rp\n' | cat -v
car^Ht
bike^Mp

# NUL character
$ printf 'car\0jeep\0bus\0' | cat -v
car^@jeep^@bus^@

# form-feed and vertical-tab
$ printf '1 2\t3\f4\v5\n' | cat -v
1 2      3^L4^K5
```

The `-v` option doesn't cover the newline and tab characters. You can use the `-T` option to spot tab characters.

```
$ printf 'good food\tnice dice\n' | cat -T
good food^Inice dice
```

The `-E` option adds a `$` marker at the end of input lines. This is useful to spot invisible trailing characters.

```
$ printf 'ice  \nwater\n cool  \n' | cat -E
ice  $
water$
cool  $
```

The following options combine two or more of the above options:

- `-e` option is equivalent to `-vE`
- `-t` option is equivalent to `-vT`
- `-A` option is equivalent to `-vET`

```
$ printf 'car\bt\nbike\rp\n' | cat -e
car^Ht$
bike^Mp$

$ printf '1 2\t3\f4\v5\n' | cat -t
1 2^I3^L4^K5

$ printf '1 2\t3\f4\v5\n' | cat -A
1 2^I3^L4^K5$
```

Useless use of cat

Using `cat` to view the contents of a file, to concatenate them, etc is well and good. But, using `cat` when it is not needed is a bad habit that you should avoid. See [wikipedia: UUOC](#) and [Useless Use of Cat Award](#) for more details.

Most commands that you'll see in this book can directly work with file arguments, so you shouldn't use `cat` and pipe the contents for such cases. Here's a single file example:

```
# useless use of cat
$ cat greeting.txt | sed -E 's/\w+/\L\u&/g'
Hi There
Have A Nice Day

# sed can handle file arguments
$ sed -E 's/\w+/\L\u&/g' greeting.txt
Hi There
Have A Nice Day
```

If you prefer having the file argument before the command, you can still use your shell's redirection feature to supply input data instead of `cat`. This also applies to commands like `tr` that do not accept file arguments.

```
# useless use of cat
$ cat greeting.txt | tr 'a-z' 'A-Z'
HI THERE
HAVE A NICE DAY

# use shell redirection instead
$ <greeting.txt tr 'a-z' 'A-Z'
HI THERE
HAVE A NICE DAY
```

Such useless use of `cat` might not have a noticeable negative impact unless you are dealing with large input files. Especially for commands like `tac` and `tail` which will have to wait for all the data to be read instead of directly processing from the end of the file if they had been passed as arguments (or using shell redirection).

If you are dealing with multiple files, then the use of `cat` will depend upon the results desired. Here's some examples:

```
# match lines containing 'o' or '0'
# -n option adds line number prefix
$ cat greeting.txt fruits.txt nums.txt | grep -n '[o0]'
5:mango
8:1000
$ grep -n '[o0]' greeting.txt fruits.txt nums.txt
fruits.txt:3:mango
nums.txt:3:1000

# count the number of lines containing 'o' or '0'
$ grep -c '[o0]' greeting.txt fruits.txt nums.txt
greeting.txt:0
fruits.txt:1
nums.txt:1
$ cat greeting.txt fruits.txt nums.txt | grep -c '[o0]'
2
```

For some use cases like in-place editing with `sed`, you can't use `cat` or shell redirection at all. The files have to be passed as arguments only. To conclude, don't use `cat` just to pass the input as `stdin` for another command unless you really need to.

tac

`tac` will display the input lines in reversed order. If you pass multiple input files, each file content will be reversed separately. Here's some examples:

```
# won't be same as: cat greeting.txt fruits.txt | tac
$ tac greeting.txt fruits.txt
Have a nice day
Hi there
mango
papaya
banana

$ printf 'apple\nbanana\ncherry\n' | tac
cherry
banana
apple
```



If the last line of input doesn't end with a newline, the output will also not have that newline character.

```
$ printf 'apple\nbanana\ncherry' | tac
cherrybanana
apple
```

Reversing input lines makes some of the text processing tasks easier. For example, if there are multiple matches but you want only the last such match. See my ebooks on [GNU sed](#) and [GNU awk](#) for more such use cases.

```
$ cat log.txt
--> warning 1
a,b,c,d
42
--> warning 2
x,y,z
--> warning 3
4,3,1

$ tac log.txt | grep -ml 'warning'
--> warning 3

$ tac log.txt | sed '/warning/q' | tac
--> warning 3
4,3,1
```

The `log.txt` input file has multiple lines containing `warning`. The task is to fetch lines based on the last match. Tools like `grep` and `sed` have features to easily match the first occurrence, so applying `tac` on the input helps to reverse the condition from last match to first match. Another benefit is that the first `tac` will stop reading input contents after the match is found in the above examples.



Use the `rev` command if you want each input line to be reversed character wise.

Customize line separator for tac

By default, the newline character is used to split the input content into *lines*. You can use the `-s` option to specify a different string to be used as the separator.

```
# use NUL as the line separator
# -s '$'\0' can also be used instead of -s '' if ANSI-C quoting is supported
$ printf 'car\0jeep\0bus\0' | tac -s '' | cat -v
bus^@jeep^@car^@

# as seen before, last entry should also have the separator
# otherwise it won't be present in the output
$ printf 'apple banana cherry' | tac -s ' ' | cat -e
cherrybanana apple $
$ printf 'apple banana cherry ' | tac -s ' ' | cat -e
cherry banana apple $
```

When the custom separator occurs before the content of interest, use the `-b` option to print

those separators before the content in the output as well.

```
$ cat body_sep.txt
%=%=
apple
banana
%=%=
red
green

$ tac -b -s '%=%= ' body_sep.txt
%=%=
red
green
%=%=
apple
banana
```

The separator will be treated as a regular expression if you use the `-r` option as well.

```
$ cat shopping.txt
apple 50
toys 5
Pizza 2
mango 25
Banana 10

# separator character is 'a' or 'm' at the start of a line
$ tac -b -rs '^[am]' shopping.txt
mango 25
Banana 10
apple 50
toys 5
Pizza 2

# alternate solution for: tac log.txt | sed '/warning/q' | tac
# separator is zero or more characters from the start of a line till 'warning'
$ tac -b -rs '^. *warning' log.txt | awk '/warning/ && ++c==2{exit} 1'
--> warning 3
4,3,1
```



See [Regular Expressions](#) chapter from my **GNU grep** ebook if you want to learn about regexp syntax and features.

head and tail

`cat` is useful to view entire contents of file(s). Pagers like `less` can be used if you are working with large files (`man` pages for example). Sometimes though, you just want a peek at the starting or ending lines of input files. Or, you know the line numbers for the information you are looking for. In such cases, you can use `head` or `tail` or a combination of both these commands to extract the content you want.

Leading and trailing lines

Consider this sample file, with line numbers prefixed for convenience.

```
$ cat sample.txt
1) Hello World
2)
3) Hi there
4) How are you
5)
6) Just do-it
7) Believe it
8)
9) banana
10) papaya
11) mango
12)
13) Much ado about nothing
14) He he he
15) Adios amigo
```

By default, `head` and `tail` will display the first and last 10 lines respectively.

```
$ head sample.txt
1) Hello World
2)
3) Hi there
4) How are you
5)
6) Just do-it
7) Believe it
8)
9) banana
10) papaya

$ tail sample.txt
6) Just do-it
7) Believe it
8)
9) banana
10) papaya
11) mango
12)
13) Much ado about nothing
```

```
14) He he he
15) Adios amigo
```

If there are less than 10 lines in the input, only those lines will be displayed.

```
# seq command will be discussed in detail later, generates 1 to 4 here
# same as: seq 4 | tail
$ seq 4 | head
1
2
3
4
```

You can use the `-nN` option to customize the number of lines (`N`) needed.

```
# first three lines
# space between -n and N is optional
$ head -n3 sample.txt
1) Hello World
2)
3) Hi there

# last two lines
$ tail -n2 sample.txt
14) He he he
15) Adios amigo
```

Excluding the last N lines

By using `head -n -N`, you can get all the input lines except the ones you'll get when you use the `tail -nN` command.

```
# except the last 11 lines
# space between -n and -N is optional
$ head -n -11 sample.txt
1) Hello World
2)
3) Hi there
4) How are you
```

Starting from Nth line

By using `tail -n +N`, you can get all the input lines except the ones you'll get when you use the `head -n(N-1)` command.

```
# all lines starting from the 11th line
# space between -n and +N is optional
$ tail -n +11 sample.txt
11) mango
12)
13) Much ado about nothing
14) He he he
```

15) Adios amigo

Multiple input files

If you pass multiple input files to the `head` and `tail` commands, each file will be processed separately. By default, the output is nicely formatted with filename headers and empty line separators.

```
$ seq 2 | head -n1 greeting.txt -  
==> greeting.txt <==  
Hi there  
  
==> standard input <==  
1
```

You can use the `-q` option to avoid filename headers and empty line separators.

```
$ tail -q -n2 sample.txt nums.txt  
14) He he he  
15) Adios amigo  
42  
1000
```

Byte selection

The `-c` option works similar to the `-n` option, but with bytes instead of lines. In the below examples, newline characters have been added to the output for illustration purposes.

```
# first three characters  
$ printf 'apple pie' | head -c3  
app  
  
# last three characters  
$ printf 'apple pie' | tail -c3  
pie  
  
# excluding last four characters  
$ printf 'car\njeep\nbus\n' | head -c -4  
car  
jeep  
  
# all characters starting from fifth character  
$ printf 'car\njeep\nbus\n' | tail -c +5  
jeep  
bus
```

Since `-c` works byte wise, it may not be suitable for multibyte characters:

```
# all input characters in this example occupy two bytes each  
$ printf 'αλεπού' | head -c2  
α  
  
# ð occupies three bytes
```

```
$ printf 'cağē' | tail -c4  
ğē
```

Range of lines

You can select a range of lines by combining both `head` and `tail` commands.

```
# 9th to 11th lines  
# same as: head -n11 sample.txt | tail -n3  
$ tail -n +9 sample.txt | head -n3  
9) banana  
10) papaya  
11) mango  
  
# 6th to 7th lines  
# same as: tail -n +6 sample.txt | head -n2  
$ head -n7 sample.txt | tail -n2  
6) Just do-it  
7) Believe it
```



See [unix.stackexchange: line X to line Y on a huge file](#) for performance comparison with other commands like `sed`, `awk`, etc.

NUL separator

The `-z` option sets the NUL character as the line separator instead of the newline character.

```
$ printf 'car\0jeep\0bus\0' | head -z -n2 | cat -v  
car^@jeep^@  
  
$ printf 'car\0jeep\0bus\0' | tail -z -n2 | cat -v  
jeep^@bus^@
```

Further Reading

- [wikipedia: File monitoring with tail -f and -F options](#)
- [unix.stackexchange: How does the tail -f option work?](#)
- [How to deal with output buffering?](#)

tr

`tr` helps you to map one set of characters to another set of characters. Features like range, repeats, character sets, squeeze, complement, etc makes it a must know text processing tool.

To be precise, `tr` can handle only bytes. Multibyte character processing isn't supported yet.

Translation

Here's some examples that map one set of characters to another. As a good practice, always enclose the sets in single quotes to avoid issues due to shell metacharacters.

```
# 'l' maps to '1', 'e' to '3', 't' to '7' and 's' to '5'
$ echo 'leet speak' | tr 'lets' '1375'
1337 5p3ak

# example with shell metacharacters
$ echo 'apple;banana;cherry' | tr ; :
tr: missing operand
Try 'tr --help' for more information.
$ echo 'apple;banana;cherry' | tr ';' ':'
apple:banana:cherry
```

You can use `-` between two characters to construct a range (ascending order only).

```
# uppercase to lowercase
$ echo 'HELLO WORLD' | tr 'A-Z' 'a-z'
hello world

# swap case
$ echo 'Hello World' | tr 'a-zA-Z' 'A-Za-z'
hELLO wORLD

# rot13
$ echo 'Hello World' | tr 'a-zA-Z' 'n-za-mN-ZA-M'
Uryyb Jbeyq
$ echo 'Uryyb Jbeyq' | tr 'a-zA-Z' 'n-za-mN-ZA-M'
Hello World
```

`tr` works only on `stdin` data, so use shell input redirection for file input.

```
$ tr 'a-z' 'A-Z' <greeting.txt
HI THERE
HAVE A NICE DAY
```

Different length sets

If the second set is longer, the extra characters are simply ignored. If the first set is longer, the last character of the second set is reused for the missing mappings.

```
# only abc gets converted to uppercase
$ echo 'apple banana cherry' | tr 'abc' 'A-Z'
Apple BAnAnA Cherry
```

```
# c-z will be converted to C
$ echo 'apple banana cherry' | tr 'a-z' 'ABC'
ACCCC BACACA CCCCC
```

You can use the `-t` option to truncate the first set so that it matches the length of the second set.

```
# d-z won't be converted
$ echo 'apple banana cherry' | tr -t 'a-z' 'ABC'
Apple BAnAnA Cherry
```

You can also use `[c*n]` notation to repeat a character `c` by `n` times. You can specify `n` in decimal format or octal format (starts with `0`). If `n` is omitted, the character `c` is repeated as many times as needed to equalize the length of the sets.

```
# a-e will be translated to A
# f-z will be uppercased
$ echo 'apple banana cherry' | tr 'a-z' '[A*5]F-Z'
APPLA AANANA AHARRY

# a-c and x-z will be uppercased
# rest of the characters will be translated to -
$ echo 'apple banana cherry' | tr 'a-z' 'ABC[-*]XYZ'
A---- BA-A-A C----Y
```

Escape sequences and character sets

Certain characters like newline, tab, etc can be represented using escape sequences. You can also specify characters using `\NNN` octal representation.

```
# same as: tr '\011' '\072'
$ printf 'apple\tbanana\tcherry\n' | tr '\t' ':'
apple:banana:cherry

$ echo 'apple:banana:cherry' | tr ':' '\n'
apple
banana
cherry
```

Certain commonly useful groups of characters like alphabets, digits, punctuation, etc have named character sets that you can use instead of manually creating the sets. Only `[:lower:]` and `[:upper:]` can be used by default, others will require `-d` or `-s` options.

```
# same as: tr 'a-z' 'A-Z' <greeting.txt
$ tr '[:lower:]' '[:upper:]' <greeting.txt
HI THERE
HAVE A NICE DAY
```

To override the special meaning for `-` and `\` characters, you can escape them using the `\` character. You can also place the `-` character at the end of a set to represent it literally. Can you reason out why placing the `-` character at the start of a set can cause issues?

```
$ echo '/python-projects/programs' | tr '/-' '\\_'
\python_projects\programs
```



See [tr manual](#) for more details and a list of all the escape sequences and character sets.

Deleting characters

Use the `-d` option to specify a set of characters to be deleted.

```
$ echo '2021-08-12' | tr -d '-'
20210812

$ s='"Hi", there! How *are* you? All fine here.'
$ echo "$s" | tr -d '[:punct:]'
Hi there How are you All fine here
```

Complement

The `-c` option will invert the first set of characters. This is often used in combination with the `-d` option.

```
$ s='"Hi", there! How *are* you? All fine here.'

# retain alphabets, whitespaces, period, exclamation and question mark
$ echo "$s" | tr -cd 'a-zA-Z.![:space:]'
Hi there! How are you? All fine here.
```

If you use `-c` for translation, you can only provide a single character for the second set. In other words, all the characters except those provided by the first set will be mapped to the character specified by the second set.

```
$ s='"Hi", there! How *are* you? All fine here.'

$ echo "$s" | tr -c 'a-zA-Z.![:space:]' '1%'
tr: when translating with complemented character classes,
string2 must map all characters in the domain to one

$ echo "$s" | tr -c 'a-zA-Z.![:space:]' '%'
%Hi%% there! How %are% you? All fine here.
```

Squeeze

The `-s` option changes consecutive repeated characters to a single copy of that character.

```
# squeeze lowercase alphabets
$ echo 'hhoowww aaaaaareeeeeee yyuuuuu!!' | tr -s 'a-z'
how are you!!

# translate and squeeze
$ echo 'hhoowww aaaaaareeeeeee yyuuuuu!!' | tr -s 'a-z' 'A-Z'
HOW ARE YOU!!

# delete and squeeze
```

```
$ echo 'hhoowww aaaaaareeeee yyouuuu!!' | tr -sd '!' 'a-z'
how are you

# squeeze other than lowercase alphabets
$ echo 'how are you!!!!' | tr -cs 'a-z'
how are you!
```


cut

`cut` is a handy tool for many field processing use cases. The features are limited compared to `awk` and `perl` commands, but the reduced scope also leads to faster processing.

Individual field selections

By default, `cut` splits the input content into fields based on the tab character. You can use the `-f` option to select a desired field from each input line. To extract multiple fields, specify the selections separated by the comma character.

```
# second field
$ printf 'apple\tbanana\tcherry\n' | cut -f2
banana

# first and third field
$ printf 'apple\tbanana\tcherry\n' | cut -f1,3
apple  cherry
```

`cut` will always display the selected fields in ascending order. Field duplication will be ignored as well.

```
# same as: cut -f1,3
$ printf 'apple\tbanana\tcherry\n' | cut -f3,1
apple  cherry

# same as: cut -f1,2
$ printf 'apple\tbanana\tcherry\n' | cut -f1,1,2,1,2,1,1,2
apple  banana
```

By default, `cut` uses the newline character as the line separator. `cut` will add a newline character to the output even if the last input line doesn't end with a newline.

```
$ printf 'good\tfood\ntip\ttap' | cut -f2
food
tap
```

Field ranges

You can use the `-` character to specify field ranges. You can skip the starting or ending range, but not both.

```
# 2nd, 3rd and 4th fields
$ printf 'apple\tbanana\tcherry\tdates\n' | cut -f2-4
banana  cherry  dates

# all fields from the start till the 3rd field
$ printf 'apple\tbanana\tcherry\tdates\n' | cut -f-3
apple  banana  cherry

# all fields from the 3rd field till the end
$ printf 'apple\tbanana\tcherry\tdates\n' | cut -f3-
cherry  dates
```

Input field delimiter

Use the `-d` option to change the input delimiter. Only a single byte character is allowed. By default, the output delimiter will be same as the input delimiter.

```
$ cat scores.csv
Name,Maths,Physics,Chemistry
Ith,100,100,100
Cy,97,98,95
Lin,78,83,80

$ cut -d, -f2,4 scores.csv
Maths,Chemistry
100,100
97,95
78,80

# use quotes if the delimiter is a shell metacharacter
$ echo 'one;two;three;four' | cut -d; -f3
cut: option requires an argument -- 'd'
Try 'cut --help' for more information.
-f3: command not found
$ echo 'one;two;three;four' | cut -d';' -f3
three
```

Output field delimiter

Use the `--output-delimiter` option to customize the output separator to any string of your choice. The string is treated literally. Depending on your shell you can use [ANSI-C quoting](#) to allow escape sequences.

```
# same as: tr '\t' ','
$ printf 'apple\tbanana\tcherry\n' | cut --output-delimiter=, -f1-
apple,banana,cherry

# multicharacter example
$ echo 'one;two;three;four' | cut -d';' --output-delimiter=' : ' -f1,3-
one : three : four

# ANSI-C quoting example
# depending on your environment, you can also press Ctrl+v and then Tab key
$ echo 'one;two;three;four' | cut -d';' --output-delimiter='${t}' -f1,3-
one      three      four

# newline as the output field separator
$ echo 'one;two;three;four' | cut -d';' --output-delimiter='${n}' -f2,4
two
four
```

Complement

The `--complement` option allows you to invert the field selections.

```
# except second field
$ printf 'apple ball cat\n1 2 3 4 5' | cut --complement -d' ' -f2
apple cat
1 3 4 5

# except first and third fields
$ printf 'apple ball cat\n1 2 3 4 5' | cut --complement -d' ' -f1,3
ball
2 4 5
```

Suppress lines without delimiters

By default, lines not containing the input delimiter will still be part of the output. You can use the `-s` option to suppress such lines.

```
$ cat mixed_fields.csv
1,2,3,4
hello
a,b,c

# second line doesn't have the comma separator
# by default, such lines will be part of the output
$ cut -d, -f2 mixed_fields.csv
2
hello
b

# use -s option to suppress such lines
$ cut -sd, -f2 mixed_fields.csv
2
b

$ cut --complement -sd, -f2 mixed_fields.csv
1,3,4
a,c
```

If a line contains the specified delimiter but doesn't have the field number requested, you'll get a blank line. The `-s` option has no effect on such lines.

```
$ printf 'apple ball cat\n1 2 3 4 5' | cut -d' ' -f4

4
```

Character selections

You can use the `-b` or `-c` options to select specified bytes from each input line. The syntax is same as the `-f` option. The `-c` option is intended for multibyte character selection, but for now it works exactly as the `-b` option. Character selection is useful for working with [fixed-width fields](#).

```
$ printf 'apple\tbanana\tcherry\n' | cut -c2,8,11
pan
```

```
$ printf 'apple\tbanana\tcherry\n' | cut -c2,8,11 --output-delimiter=-
p-a-n

$ printf 'apple\tbanana\tcherry\n' | cut -c-5
apple

$ printf 'apple\tbanana\tcherry\n' | cut --complement -c13-
apple  banana

$ printf 'cat-bat\ndog:fog\nget;pet' | cut -c5-
bat
fog
pet
```

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `cut` will ensure to add a final NUL character even if not present in the input.

```
$ printf 'good-food\0tip-tap\0' | cut -zd- -f2 | cat -v
food^@tap^@
```

Alternatives

Here's some alternate commands you can explore if `cut` isn't enough to solve your task.

- [hck](#) — supports regexp delimiters, field reordering, header based selection, etc
- [xsv](#) — fast CSV command line toolkit
- [rcut](#) — my `bash+awk` script, supports regexp delimiters, field reordering, negative indexing, etc
- [awk](#) — my ebook on `GNU awk` one-liners
- [perl](#) — my ebook on `perl` one-liners

seq

The `seq` command is a handy tool to generate a sequence of numbers in ascending or descending order. Both integer and floating-point numbers are supported. You can also customize the formatting for numbers and the separator between them.

Integer sequences

You need three numbers to generate an arithmetic progression — **start**, **step** and **stop**. When you pass only a single number as the stop value, the default start and step values are assumed to be `1`.

```
# start=1, step=1 and stop=3
$ seq 3
1
2
3
```

Passing two numbers are considered as start and stop values (in that order).

```
# start=25434, step=1 and stop=25437
$ seq 25434 25437
25434
25435
25436
25437

# start=-5, step=1 and stop=-3
$ seq -5 -3
-5
-4
-3
```

When you want to specify all the three numbers, the order is start, step and stop.

```
# start=1000, step=5 and stop=1010
$ seq 1000 5 1010
1000
1005
1010
```

By using a negative step value, you can generate sequences in descending order.

```
# no output
$ seq 3 1

# need to explicitly use a negative step value
$ seq 3 -1 1
3
2
1

$ seq 5 -5 -10
5
```

```
0
-5
-10
```

Floating-point sequences

Since `1` is the default start and step values, you need to change at least one of them to get floating-point sequences.

```
$ seq 0.5 3
0.5
1.5
2.5

$ seq 0.25 0.33 1.12
0.25
0.58
0.91
```

[E scientific notation](#) is also supported.

```
$ seq 1.2e2 1.22e2
120
121
122

$ seq 1.2e2 0.752 1.22e2
120.000
120.752
121.504
```

Customizing separator

You can use the `-s` option to change the separator between the numbers of a sequence. Multiple characters are allowed. Depending on your shell you can use [ANSI-C quoting](#) to use escapes like `\t` instead of a literal tab character. A newline is always added at the end of the output.

```
$ seq -s ' ' 4
1 2 3 4

$ seq -s: -2 0.75 3
-2.00:-1.25:-0.50:0.25:1.00:1.75:2.50

$ seq -s' - ' 4
1 - 2 - 3 - 4

$ seq -s$'\n\n' 4
1

2

3
```

Leading zeros

By default, the output will not have leading zeros, even if they are part of the numbers passed to the command.

```
$ seq 008 010
8
9
10
```

The `-w` option will equalize the width of the output numbers using leading zeros. The largest width between the start and stop values will be used.

```
$ seq -w 8 10
08
09
10

$ seq -w 0003
0001
0002
0003
```

printf style formatting

You can use the `-f` option for `printf` style floating-point number formatting. See [bash manual: printf](#) for more details on formatting options.

```
$ seq -f'%g' -s: 1 0.75 3
1:1.75:2.5

$ seq -f'%0.4f' -s: 1 0.75 3
1.0000:1.7500:2.5000

$ seq -f'%0.3e' 1.2e2 0.752 1.22e2
1.200e+02
1.208e+02
1.215e+02
```

Limitations

As per the manual:

On most systems, `seq` can produce whole-number output for values up to at least 2^{53} . Larger integers are approximated. The details differ depending on your floating-point implementation.

```
# example with approximate values
$ seq 1000000000000000000000 3 1000000000000000000010
```

```
10000000000000000000000000
10000000000000000000000000
10000000000000000000000008
10000000000000000000000008
```

However, note that when limited to non-negative whole numbers, an increment of `1` and no format-specifying option, `seq` can print arbitrarily large numbers.

[illegible]

shuf

The `shuf` command helps you randomize input lines. And there are features to limit the number of output lines, repeat lines and even generate random positive integers.

Randomize input lines

By default, `shuf` will randomize the order of input lines. Here's an example:

```
$ cat purchases.txt
coffee
tea
washing powder
coffee
toothpaste
tea
soap
tea

$ shuf purchases.txt
tea
coffee
tea
toothpaste
soap
coffee
washing powder
tea
```



You can use the `--random-source=FILE` option to provide your own source for randomness. With this option, the output will be the same across multiple runs. See [Sources of random data](#) for more details.



`shuf` doesn't accept multiple input files. Use `cat` for such cases.

Limit output lines

Use the `-n` option to limit the number of lines you want in the output. If the value is greater than the number of lines in the input, it would be similar to not using the `-n` option.

```
$ printf 'apple\nbanana\ncherry' | shuf -n2
cherry
apple
```



As seen in the example above, `shuf` will add a newline character if it is not present for the last input line.

Repeated lines

The `-r` option helps if you want to allow input lines to be repeated. This option is usually paired with `-n` to limit the number of lines in the output.

```
$ cat fruits.txt
banana
papaya
mango

$ shuf -n3 -r fruits.txt
banana
mango
banana

$ shuf -n5 -r fruits.txt
papaya
banana
mango
papaya
papaya
```



If a limit using `-n` is not specified, `shuf -r` will produce output lines indefinitely.

Specify input as arguments

You can use the `-e` option to specify multiple input lines as arguments to the command.

```
# quote the arguments as necessary
$ shuf -e hi there 'hello world' good
hello world
good
hi
there

$ shuf -n1 -e red green blue
blue

$ shuf -n4 -r -e red green blue
blue
green
red
blue
```

The shell will autocomplete unquoted glob patterns (provided there are files that match the given expression). You can thus easily construct a solution to get a random selection of files matching the given glob pattern.

```
$ echo *.csv
marks.csv mixed_fields.csv report_1.csv report_2.csv scores.csv
```

```
$ shuf -n2 -e *.csv  
scores.csv  
marks.csv
```

Generate random numbers

The `-i` option will help you generate random positive integers.

```
$ shuf -i 5-8
5
8
7
6

$ shuf -n3 -i 100-200
170
112
148

$ shuf -n5 -r -i 0-1
1
0
0
1
1
```



$2^{64} - 1$ is the maximum allowed integer when I tested it on my machine.

```
$ shuf -i 18446744073709551612-18446744073709551615
18446744073709551615
18446744073709551614
18446744073709551612
18446744073709551613
```

```
$ shuf -i 18446744073709551612-18446744073709551616
shuf: invalid input range: '18446744073709551616':
Value too large for defined data type
```

[illegible]

seq can also help when you need negative and floating-point numbers.

```
$ seq -10 -8 | shuf
-9
-10
-8
```

```
$ seq -f'%.4f' 100 0.25 3000 | shuf -n3
1627.7500
1303.5000
2466.2500
```



See [unix.stackexchange: generate random strings](https://unix.stackexchange.com/questions/111111/generate-random-strings) if numbers aren't enough for you.

Specifying output file

The `-o` option can be used to specify the output file to be used for saving the results. You can use this for in-place editing as well, since `shuf` reads the entire input before opening the output file.

```
$ shuf nums.txt -o rand_nums.txt

$ cat rand_nums.txt
42
1000
3.14
```

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `shuf` will ensure to add a final NUL character even if not present in the input.

```
$ printf 'apple\0banana\0cherry' | shuf -z -n2 | cat -v
cherry^@banana^@
```

paste

`paste` is typically used to merge two or more files column wise. It also has a handy feature for serializing data.

Concatenating files column wise

Consider these two input files:

```
$ cat colors_1.txt
Blue
Brown
Orange
Purple
Red
Teal
White
$ cat colors_2.txt
Black
Blue
Green
Orange
Pink
Red
White
```

By default, `paste` adds a tab character between corresponding lines of input files.

```
$ paste colors_1.txt colors_2.txt
Blue      Black
Brown     Blue
Orange    Green
Purple     Orange
Red       Pink
Teal      Red
White     White
```

You can use the `-d` option to change the delimiter between the columns. The separator is added even if the data has been exhausted for some of the input files. Here's some examples with single character delimiters, multicharacter separation will be discussed later.

```
$ seq 5 | paste -d, - <(seq 6 10)
1,6
2,7
3,8
4,9
5,10

# quote the delimiter if it is a shell metacharacter
$ paste -d'|' <(seq 3) <(seq 4 5) <(seq 6 8)
1|4|6
2|5|7
3||8
```

Use empty string if you don't want any delimiter between the columns. You can also use `\0` for this case, but that'd be confusing since it is typically used to mean the NUL character.

```
# note that the space between -d and empty string is necessary here
$ paste -d ' ' <(seq 3) <(seq 6 8)
16
27
38
```



You can pass the same filename multiple times too, they will be treated as if they are separate inputs. This doesn't apply for `stdin` though, which is a special case as discussed in a later section.

Interleaving lines

By setting the newline character as the delimiter, you'll get interleaved lines.

```
$ paste -d '\n' <(seq 11 13) <(seq 101 103)
11
101
12
102
13
103
```

Multiple columns from single input

If you use `-` multiple times, `paste` will consume a line from `stdin` data every time `-` is encountered. This is different from using the same filename multiple times, in which case they are treated as separate inputs.

This special case for `stdin` data is useful to combine consecutive lines using the given delimiter. Here's some examples to help you understand this feature better:

```
# two columns
$ seq 10 | paste -d, - -
1,2
3,4
5,6
7,8
9,10

# five columns
$ seq 10 | paste -d: - - - - -
1:2:3:4:5
6:7:8:9:10

# use input redirection for file input
$ <colors_1.txt paste -d: - - -
Blue:Brown:Orange
Purple:Red:Teal
White::
```

Here's an example with both `stdin` and file arguments:

```
$ seq 6 | paste - nums.txt -
1      3.14    2
3      42      4
5      1000    6
```

If you don't want to manually type the number of `-` required, you can use this `printf` trick:

```
# the string before %.s is repeated based on the number of arguments
$ printf 'x %.s' a b c
x x x
$ printf -- '- %.s' {1..5}
- - - - -

$ seq 10 | paste -d, $(printf -- '- %.s' {1..5})
1,2,3,4,5
6,7,8,9,10
```



See [this stackoverflow thread](#) for more details about the `printf` solution and other alternatives.

Multicharacter delimiters

The `-d` option accepts a list of characters (bytes to be precise) to be used one by one between the different columns. If the number of characters is less than the number of separators required, the characters are reused from the beginning and this cycle repeats until all the columns are done. If the number of characters is greater than the number of separators required, the extra characters are simply discarded.

```
# , is used between 1st and 2nd column
# - is used between 2nd and 3rd column
$ paste -d',-' <(seq 3) <(seq 4 6) <(seq 7 9)
1,4-7
2,5-8
3,6-9

# only 3 separators are needed, the rest are discarded
$ paste -d',-:;.[]' <(seq 3) <(seq 4 6) <(seq 7 9) <(seq 10 12)
1,4-7:10
2,5-8:11
3,6-9:12

# 2 characters given, 4 separators needed
# paste will reuse from the start of the list
$ seq 10 | paste -d':,' - - - -
1:2,3:4,5
6:7,8:9,10
```

You can use empty files to get multicharacter separation between the columns.

```
$ paste -d' : ' <(seq 3) /dev/null /dev/null <(seq 4 6)
1 : 4
2 : 5
3 : 6

# create empty file to avoid typing /dev/null too many times
$ > e
$ paste -d' : - ' <(seq 3) e e <(seq 4 6) e e <(seq 7 9)
1 : 4 - 7
2 : 5 - 8
3 : 6 - 9
```

Serialize

The `-s` option allows you to combine all the input lines from a file into a single line using the given delimiter. `paste` will ensure to add a final newline character even if it isn't present in the input.

```
# <colors_1.txt tr '\n' ',' will give you a trailing comma
# paste changes the separator between the lines only
$ paste -sd, colors_1.txt
Blue,Brown,Orange,Purple,Red,Teal,White

# newline gets added at the end even if not present in the input
$ printf 'apple\nbanana\ncherry' | paste -sd-
apple-banana-cherry
```

If multiple files are passed, serialization of each file is displayed on separate lines.

```
$ paste -sd: colors_1.txt colors_2.txt
Blue:Brown:Orange:Purple:Red:Teal:White
Black:Blue:Green:Orange:Pink:Red:White

$ paste -sd, <(seq 3) <(seq 5 9)
1,2,3
5,6,7,8,9
```

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `paste` will ensure to add a final NUL character even if not present in the input.

```
$ printf 'a\0b\0c\0d\0' | paste -z -d: - - | cat -v
a:b^@c:d^@
```


pr

Paginate or columnate FILE(s) for printing.

As stated in the above quote from the manual, the `pr` command is mainly used for those two tasks. This book will discuss only the columnate features and some miscellaneous tasks.

Here's a pagination example if you are interested in exploring further. The `pr` command will add blank lines, a header and so on to make it suitable for printing.

```
$ pr greeting.txt | head
```

```
2021-08-05 14:10          greeting.txt          Page 1

Hi there
Have a nice day
```

Columnate

The `--columns` and `-a` options can be used to merge the input lines in two different ways:

- split the input file and then merge them as columns
- merge consecutive lines, similar to the `paste` command

Here's an example to get started. Note that `-N` is same as using `--columns=N` where `N` is the number of columns you want in the output. The default page width is `72`, which means each column can only have a maximum of `72/N` characters (including the separator). Tab and space characters will be used to fill the columns as needed. You can use the `-J` option to prevent `pr` from truncating longer columns. The `-t` option is used here to turn off the pagination features.

```
# split input into three parts
# each column width is 72/3 = 24 characters max
$ seq 9 | pr -3t
1          4          7
2          5          8
3          6          9
```

You can customize the separator using the `-s` option. The default is a tab character which you can change to any other string value. The `-s` option also turns off line truncation, so `-J` option isn't needed. However, the default page width of `72` can still cause issues, which will be discussed later.

```
# tab separator
$ seq 9 | pr -3ts
1      4      7
2      5      8
3      6      9
```

```
# comma separator
$ seq 9 | pr -3ts,
1,4,7
2,5,8
3,6,9

# multicharacter separator
$ seq 9 | pr -3ts' : '
1 : 4 : 7
2 : 5 : 8
3 : 6 : 9
```

Use the `-a` option to merge consecutive lines, similar to the `paste` command. One advantage is that the `-s` option supports a string value, whereas with `paste` you'd need to use workarounds to get multicharacter separation.

```
# four consecutive lines are merged
# same as: paste -d: - - - -
$ seq 8 | pr -4ats:
1:2:3:4
5:6:7:8
```

There are other differences between the `pr` and `paste` commands as well. Unlike `paste`, the `pr` command doesn't add the separator if the last row doesn't have enough columns. Another difference is that `pr` doesn't support an option to use the NUL character as the line separator.

```
$ seq 10 | pr -4ats,
1,2,3,4
5,6,7,8
9,10

$ seq 10 | paste -d, - - - -
1,2,3,4
5,6,7,8
9,10,,
```

Customizing page width

As mentioned before, the default page width is `72`. This can cause lines to be truncated, unless `-s` or `-J` options are used. There's another issue you might run into, for example:

```
$ seq 100 | pr -50ats,
pr: page width too narrow
```

$(N-1) * \text{length}(\text{separator}) + N$ is the minimum page width you need, where `N` is the number of columns required. So, for `50` columns and a separator of length `1`, you'll need a minimum width of `99`. This calculation doesn't make any assumption about the size of input lines, so you may need `-J` to ensure input lines aren't truncated.

You can use the `-w` option to change the page width. The `-w` option overrides the effect of `-s` option on line truncation, so use `-J` option as well unless you really need truncation. If

truncation is active, maximum column width is $(\text{PageWidth} - (N-1) * \text{length}(\text{separator})) / N$ rounded down to an integer value. Here's some examples:

```
# minimum width needed is 3 for N=2 and length=1
# maximum column width: (6 - 1) / 2 = 2
$ pr -w6 -2ts, greeting.txt
Hi,Ha
# use -J to avoid truncation
$ pr -J -w6 -2ts, greeting.txt
Hi there,Have a nice day

# N=3 and length=4, so minimum width needed is (3-1)*4 + 3 = 11
$ seq 6 | pr -J -w10 -3ats'::::'
pr: page width too narrow
$ seq 6 | pr -J -w11 -3ats'::::'
1:::2:::3
4:::5:::6

# you can also use a large number to avoid having to calculate the width
$ seq 6 | pr -J -w500 -3ats'::::'
1:::2:::3
4:::5:::6
```

Concatenating files column wise

Two or more input files can be merged column wise using the `-m` option. As seen before, `-t` is needed to ignore pagination features and `-s` can be used to customize the separator.

```
# same as: paste colors_1.txt colors_2.txt
$ pr -mts colors_1.txt colors_2.txt
Blue    Black
Brown   Blue
Orange  Green
Purple  Orange
Red      Pink
Teal     Red
White    White

# same as: paste -d' : ' <(seq 3) /dev/null /dev/null <(seq 4 6)
$ pr -mts' : ' <(seq 3) <(seq 4 6)
1 : 4
2 : 5
3 : 6
```

You can prefix the output with line numbers using the `-n` option. By default, this option supports up to `5` digit numbers and uses the tab character to separate the numbering and line contents. You can optionally pass two arguments to this option — maximum number of digits and the separator character. If both arguments are used, the separator should be specified first. If you want to customize the starting line number, use the `-N` option as well.

```
# maximum of 1 digit for numbering
# use : as the separator between line number and line contents
```

```
$ pr -n:1 -mts, colors_1.txt colors_2.txt
1:Blue,Black
2:Brown,Blue
3:Orange,Green
4:Purple,Orange
5:Red,Pink
6:Teal,Red
7:White,White
```

The string passed to `-s` is treated literally. Depending on your shell you can use [ANSI-C quoting](#) to allow escape sequences. Unlike `columnate`, the separator is added even if the data is missing for some of the files.

```
# greeting.txt has 2 lines
# fruits.txt has 3 lines
# same as: paste -d$'\n' greeting.txt fruits.txt
$ pr -mts$'\n' greeting.txt fruits.txt
Hi there
banana
Have a nice day
papaya

mango
```

Miscellaneous

You can use the `-d` option to double space the input contents. That is, every newline character is doubled.

```
$ pr -dt fruits.txt
banana

papaya

mango
```

The `-v` option will convert non-printing characters like carriage return, backspace, etc to their octal representations (`\NNN`).

```
$ printf 'car\b\t\r\nbike\0p\r\n' | pr -vt
car\010t\015
bike\000p\015
```

`pr -t` is a roundabout way of concatenating input files. But one advantage is that this will add a newline character at the end if not present in the input.

```
# 'cat' will not add a newline character
# so, use 'pr' if newline is needed at the end
$ printf 'a\nb\nc' | pr -t
a
b
c
```

fold and fmt

These two commands are useful to split and join lines to meet a specific line length requirement. `fmt` is smarter and usually the tool you want, but `fold` can be handy for some cases.

fold

By default, `fold` will wrap lines that are greater than 80 bytes long, which can be customized using the `-w` option. The newline character isn't part of this line length calculation. You might wonder if there are tasks where wrapping without context could be useful. One use case I can think of is the [FASTA format](#).

```
$ cat greeting.txt
Hi there
Have a nice day

# splits second line since it is greater than 10 bytes
$ fold -w10 greeting.txt
Hi there
Have a nic
e day
```

The `-s` option looks for the presence of spaces to determine the line splitting. This check is performed within the limits of the wrap length.

```
$ fold -s -w10 greeting.txt
Hi there
Have a
nice day
```

However, the `-s` option can still split words if there's no blank space before the specified width. Use `fmt` if you don't want this behavior.

```
$ echo 'hi there' | fold -s -w4
hi
ther
e
```

The `-b` option will cause `fold` to treat tab, backspace, and carriage return characters as if they were a single byte character.

```
# tab can occupy up to 8 columns
$ printf 'a\tb\tc\tl\t2\t3\n' | fold -w6
a
b
c
1
2
3
```

```
# here, tab will be treated as if it occupies only 1 column
$ printf 'a\tb\tc\tl\t2\t3\n' | fold -b -w6
a      b      c
1      2      3
```

fmt

The `fmt` command makes a smarter decision based on sentences, paragraphs and other details. Here's an example that splits a single line (taken from the documentation of `fmt` command) into several lines. The default formatting is 93% of 75 columns. The `-w` option controls the width parameter and the `-g` option controls the percentage of columns.

```
$ fmt info_fmt.txt
fmt prefers breaking lines at the end of a sentence, and tries to
avoid line breaks after the first word of a sentence or before the last
word of a sentence. A sentence break is defined as either the end of a
paragraph or a word ending in any of '.?!', followed by two spaces or
end of line, ignoring any intervening parentheses or quotes. Like TeX,
fmt reads entire "paragraphs" before choosing line breaks; the algorithm
is a variant of that given by Donald E. Knuth and Michael F. Plass in
"Breaking Paragraphs Into Lines", Software-Practice & Experience 11,
11 (November 1981), 1119-1184.
```

Unlike `fold`, words are not split even if they exceed the maximum line width. Another difference is that `fmt` will add a final newline character if it isn't present in the input.

```
$ printf 'hi there' | fmt -w4
hi
there
```

The `fmt` command also allows you to join lines together that are shorter than the specified width. As mentioned before, paragraphs are taken into consideration, so empty lines will prevent merging. The `-s` option will disable line merging.

```
$ cat sample.txt
1) Hello World
2)
3) Hi there
4) How are you
5)
6) Just do-it
7) Believe it
8)
9) banana
10) papaya
11) mango
12)
13) Much ado about nothing
14) He he he
15) Adios amigo

# 'cut' here helps to ignore first 4 characters of sample.txt
```

```
$ cut -c5- sample.txt | fmt -w30
Hello World

Hi there How are you

Just do-it Believe it

banana papaya mango

Much ado about nothing He he
he Adios amigo
```

The `-u` option will change multiple spaces to a single space. Excess spacing between sentences will be changed to two spaces.

```
$ printf 'Hi   there.   Have a nice   day\n' | fmt -u
Hi there.  Have a nice day
```

There are options that control indentation, option to format only lines with a specific prefix and so on. See [fmt documentation](#) for more details.

sort

The `sort` command provides a wide variety of features. In addition to [lexicographic ordering](#), it supports various numerical formats. You can also sort based on particular column(s). And there are nifty features like merging already sorted input, debugging, whether input is sorted and so on.

Default sort and Collating order

By default, `sort` orders the input in ascending order. If you know about [ASCII codepoints](#), do you agree that the following two examples are showing the correct expected output?

```
$ cat greeting.txt
Hi there
Have a nice day
# extract and sort space separated words
$ <greeting.txt tr ' ' '\n' | sort
a
day
Have
Hi
nice
there

$ printf '(banana)\n{cherry}\n[apple]' | sort
[apple]
(banana)
{cherry}
```

From [sort manual](#):

Unless otherwise specified, all comparisons use the character collating sequence specified by the `LC_COLLATE` locale.

If you use a non-POSIX locale (e.g., by setting `LC_ALL` to `en_US`), then `sort` may produce output that is sorted differently than you're accustomed to. In that case, set the `LC_ALL` environment variable to `C`. Note that setting only `LC_COLLATE` has two problems. First, it is ineffective if `LC_ALL` is also set. Second, it has undefined behavior if `LC_CTYPE` (or `LANG`, if `LC_CTYPE` is unset) is set to an incompatible value. For example, you get undefined behavior if `LC_CTYPE` is `ja_JP.PCK` but `LC_COLLATE` is `en_US.UTF-8`.

All my `locale` settings are based on `en_IN`, which is different from the POSIX sorting order. So, the fact to remember is that `sort` obeys the rules of the current `locale`. If you want POSIX sorting, one option is to use `LC_ALL=C` as shown below.

```
$ <greeting.txt tr ' ' '\n' | LC_ALL=C sort
Have
Hi
a
day
nice
there
```



```
$ printf '(banana)\n{cherry}\n[apple]' | LC_ALL=C sort
(banana)
[apple]
{cherry}
```



Another benefit of `C` locale is that it will be significantly faster compared to Unicode parsing and sorting rules.



Use `-f` option if you want to explicitly ignore case. See also [GNU Core Utilities FAQ: Sort does not sort in normal order!](#).



See [this unix.stackexchange thread](#) if you want to create your own custom sort order.

Ignoring headers

You can use `sed -u` to consume only the header line(s) and leave the rest of the input for the `sort` command. Note that this unbuffered option is supported by `GNU sed`, might not be available with other implementations.

```
$ cat scores.csv
Name,Maths,Physics,Chemistry
Ith,100,100,100
Cy,97,98,95
Lin,78,83,80

# lq is used to quit after the first line
$ ( sed -u 'lq' ; sort ) <scores.csv
Name,Maths,Physics,Chemistry
Cy,97,98,95
Ith,100,100,100
Lin,78,83,80
```



See [this unix.stackexchange thread](#) for more ways of ignoring headers. See [bash manual: Grouping Commands](#) for more details about the `()` grouping used in the above example.

Dictionary sort

The `-d` option will consider only alphabets, numbers and blanks for sorting. Space and tab characters are considered as blanks, but this would also depend on the locale.

```
$ printf '(banana)\n{cherry}\n[apple]' | LC_ALL=C sort -d
[apple]
(banana)
{cherry}
```



Use the `-i` option if you want to ignore only the non-printing characters.

Reversed order

The `-r` option will reverse the output order. Note that this doesn't change how `sort` performs comparisons, only the output is reversed. You'll see an example later where this distinction becomes clearer.

```
$ printf 'peace\nrest\nquiet' | sort -r
rest
quiet
peace
```



In case you haven't noticed yet, `sort` adds a newline character to the final input line if it isn't present.

Numeric sort

`sort` provides various options to work with numeric formats. For most cases, the `-n` option is enough. Here's an example:

```
# lexicographic ordering isn't suited for numbers
$ printf '20\n2\n3' | sort
2
20
3

# -n helps in this case
$ printf '20\n2\n3' | sort -n
2
3
20
```

The `-n` option can handle negative and floating-point numbers as well. The decimal point and the thousands separator characters will depend on the `locale` settings.

```
$ cat mixed_numbers.txt
12,345
42
31.24
-100
42
5678

# , is the thousands separator in en_IN
# . is the decimal point in en_IN
$ sort -n mixed_numbers.txt
-100
31.24
```

```
42
42
5678
12,345
```

Use the `-g` option if your input can have the `+` prefix for positive numbers or follows the [E scientific notation](#).

```
$ cat e_notation.txt
+120
-1.53
3.14e+4
42.1e-2

$ sort -g e_notation.txt
-1.53
42.1e-2
+120
3.14e+4
```



Unless otherwise specified, `sort` will break ties by using the entire input line content. In the case of `-n` option, sorting will work even if there are extra characters after the number. Those extra characters will affect the output order if the numbers are equal. If a line doesn't start with a number (excluding blanks), it will be treated as `0`.

```
# 'b' comes before 'p'
$ printf '2 pins\n13 pens\n2 balls' | sort -n
2 balls
2 pins
13 pens

# 'z' and 'a2p' will be treated as '0'
# 'a' comes before 'z'
$ printf 'z\na2p\n13p\n2b\n-1\n    10' | sort -n
-1
a2p
z
2b
    10
13p
```

Human numeric sort

Commands like `du` (disk usage) have `-h` and `--si` options to display numbers with [SI suffixes](#) like `k`, `K`, `M`, `G` and so on. In such cases, you can use `sort -h` to order them.

```
$ cat file_size.txt
104K    power.log
316M    projects
746K    report.log
```

```
20K    sample.txt
1.4G    games

$ sort -hr file_size.txt
1.4G    games
316M    projects
746K    report.log
104K    power.log
20K     sample.txt
```

Version sort

The `-V` option is useful when you have a mix of alphabets and digits. It also helps when you want to treat digits after a decimal point as whole numbers, for example `1.10` should be greater than `1.2`.

```
$ printf '1.10\n1.2' | sort -n
1.10
1.2
$ printf '1.10\n1.2' | sort -V
1.2
1.10

$ cat versions.txt
file2
cmd5.2
file10
cmd1.6
file5
cmd5.10
$ sort -V versions.txt
cmd1.6
cmd5.2
cmd5.10
file2
file5
file10
```

Here's an example of dealing with numbers reported by the `time` command (assuming all the entries have the same format).

```
$ cat timings.txt
5m35.363s
3m20.058s
4m11.130s
3m42.833s
4m3.083s

$ sort -V timings.txt
3m20.058s
3m42.833s
```

```
4m3.083s
4m11.130s
5m35.363s
```



See [Version sort ordering](#) for more details. Note that the `ls` command uses lowercase `-v` for this task.

Random sort

The `-R` option will display the output in random order. Unlike `shuf`, this option will always place identical lines next to each other due to the implementation.

```
# the two lines with '42' will always be next to each other
# use 'shuf' if you don't want this behavior
$ sort -R mixed_numbers.txt
31.24
5678
42
42
12,345
-100
```

Unique sort

The `-u` option will keep only the first copy of lines that are deemed to be equal.

```
# (10) and [10] are deemed equal in dictionary sort
$ printf '(10)\n[20]\n[10]' | sort -du
(10)
[20]

$ cat purchases.txt
coffee
tea
washing powder
coffee
toothpaste
tea
soap
tea
$ sort -u purchases.txt
coffee
soap
tea
toothpaste
washing powder
```

As seen earlier, `-n` option will work even if there are extra characters after the number. When `-u` option is also used, only the first such copy will be retained. Use the `uniq` command if you want to remove duplicates based on the whole line.

```
$ printf '2 balls\n13 pens\n2 pins\n13 pens\n' | sort -nu
2 balls
13 pens

# note that only the output order is reversed
# use tac if you want the last duplicate to be preserved instead of first
$ printf '2 balls\n13 pens\n2 pins\n13 pens\n' | sort -r -nu
13 pens
2 balls

# use uniq when the entire line contents should be compared
$ printf '2 balls\n13 pens\n2 pins\n13 pens\n' | sort -n | uniq
2 balls
2 pins
13 pens
```

You can use the `-f` option to ignore case while determining duplicates.

```
$ printf 'cat\nbat\nCAT\ncar\nbat\n' | sort -u
bat
car
cat
CAT

# first copy between 'cat' and 'CAT' is retained
$ printf 'cat\nbat\nCAT\ncar\nbat\n' | sort -fu
bat
car
cat
```

Column sort

The `-k` option allows you to sort based on specific column(s) instead of the entire input line. By default, the empty string between non-blank and blank characters is considered as the separator and thus the blanks are also part of the field contents. The effect of blanks and mitigation will be discussed later.

The `-k` option accepts arguments in various ways. You can specify starting and ending column numbers separated by a comma. If you specify only the starting column, the last column will be used as the ending column. Usually you just want to sort by a single column, in which case the same number is specified as both the starting and ending columns. Here's an example:

```
$ cat shopping.txt
apple 50
toys 5
Pizza 2
mango 25
Banana 10

# sort based on 2nd column numbers
$ sort -k2,2n shopping.txt
Pizza 2
```

```
toys    5
Banana  10
mango   25
apple   50
```



Note that in the above example, the `-n` option was also appended to the `-k` option. This makes it specific to that column and overrides global options, if any. Also, remember that the entire line will be used to break ties, unless otherwise specified.

You can use the `-t` option to specify a single byte character as the field separator. Use `\0` to specify NUL as the separator. Depending on your shell you can use [ANSI-C quoting](#) to use escapes like `\t` instead of a literal tab character. When `-t` option is used, the field separator won't be part of the field contents.

```
# department,name,marks
$ cat marks.csv
ECE,Raj,53
ECE,Joel,72
EEE,Moi,68
CSE,Surya,81
EEE,Raj,88
CSE,Moi,62
EEE,Tia,72
ECE,Om,92
CSE,Amy,67

# name column is the primary sort key
# entire line content will be used for breaking ties
$ sort -t, -k2,2 marks.csv
CSE,Amy,67
ECE,Joel,72
CSE,Moi,62
EEE,Moi,68
ECE,Om,92
ECE,Raj,53
EEE,Raj,88
CSE,Surya,81
EEE,Tia,72
```

You can use the `-k` option multiple times to specify your own order of tie breakers. Entire line will still be used to break ties if needed.

```
# second column is the primary key
# reversed numeric sort on third column is the secondary key
# entire line will be used only if there are still tied entries
$ sort -t, -k2,2 -k3,3nr marks.csv
CSE,Amy,67
ECE,Joel,72
EEE,Moi,68
CSE,Moi,62
```

```

ECE,0m,92
EEE,Raj,88
ECE,Raj,53
CSE,Surya,81
EEE,Tia,72

# sort by month first and then the day
# -M option sorts based on abbreviated month names
$ printf 'Aug-20\nMay-5\nAug-3' | sort -t- -k1,1M -k2,2n
May-5
Aug-3
Aug-20

```

Use the `-s` option to retain the original order of input lines when two or more lines are deemed equal. You can still use multiple keys to specify your own tie breakers, `-s` only prevents the last resort comparison.

```

# -s prevents last resort comparison
# so, lines having the same value in 2nd column will retain input order
$ sort -t, -s -k2,2 marks.csv
CSE,Amy,67
ECE,Joel,72
EEE,Moi,68
CSE,Moi,62
ECE,0m,92
ECE,Raj,53
EEE,Raj,88
CSE,Surya,81
EEE,Tia,72

```

The `-u` option, as discussed earlier, will retain only the first copy of lines that are deemed equal.

```

# only the first copy of duplicates in 2nd column will be retained
$ sort -t, -u -k2,2 marks.csv
CSE,Amy,67
ECE,Joel,72
EEE,Moi,68
ECE,0m,92
ECE,Raj,53
CSE,Surya,81
EEE,Tia,72

```

Character positions within columns

The `-k` option also accepts starting and ending character positions within the columns. These are specified after the column number, separated by a `.` character. If the character position is not specified for the ending column, the last character of that column is assumed.

The character positions start with `1` for the first character. Recall that when the `-t` option is used, the field separator is not part of the field contents.


```
# based on the second column number
# 2.2 helps to ignore first character, otherwise -n won't have any effect here
$ printf 'car,(20)\njeep,[10]\ntruck,(5)\nbus,[3]' | sort -t, -k2.2,2n
bus,[3]
truck,(5)
jeep,[10]
car,(20)

# first character of the second column is the primary key
# entire line acts as the last resort tie breaker
$ printf 'car,(20)\njeep,[10]\ntruck,(5)\nbus,[3]' | sort -t, -k2.1,2.1
car,(20)
truck,(5)
bus,[3]
jeep,[10]
```

The default blanks based separation works differently. The empty string between non-blank and blank characters is considered as the separator and thus the blanks are also part of the field contents. You can use the `-b` option to ignore such leading blanks of field contents.

```
# the second column here starts with blank characters
# adjusting the character position isn't feasible due to varying blanks
$ printf 'car (20)\njeep [10]\ntruck (5)\nbus [3]' | sort -k2.2,2n
bus [3]
car (20)
jeep [10]
truck (5)

# use -b in such cases to ignore leading blanks
$ printf 'car (20)\njeep [10]\ntruck (5)\nbus [3]' | sort -k2.2b,2n
bus [3]
truck (5)
jeep [10]
car (20)
```

Debugging

The `--debug` option can help you identify issues if the output isn't what you expected. Here's the previously seen `-b` example, now with `--debug` enabled. The underscores in the debug output shows which portions of the input are used as primary key, secondary key and so on. The collating order being used is also shown in the output.

```
$ printf 'car (20)\njeep [10]\ntruck (5)\nbus [3]' | sort -k2.2,2n --debug
sort: using 'en_IN' sorting rules
sort: leading blanks are significant in key 1; consider also specifying 'b'
bus [3]
    ^ no match for key
_____
car (20)
    ^ no match for key
_____
```

```

jeep [10]
    ^ no match for key
_____
truck (5)
    ^ no match for key
_____

$ printf 'car (20)\njeep [10]\ntruck (5)\nbus [3]' | sort -k2.2b,2n --debug
sort: using 'en_IN' sorting rules
bus [3]
    -
_____
truck (5)
    -
_____
jeep [10]
    -
_____
car (20)
    -
_____

```

Check if sorted

The `-c` options helps you spot the first unsorted entry in the given input. The uppercase `-C` option is similar but only affects the exit status. Note that these options will not work for multiple inputs.

```

$ cat shopping.txt
apple  50
toys   5
Pizza  2
mango  25
Banana 10

$ sort -c shopping.txt
sort: shopping.txt:3: disorder: Pizza  2
$ echo $?
1

$ sort -C shopping.txt
$ echo $?
1

```

Specifying output file

The `-o` option can be used to specify the output file to be used for saving the results.

```

$ sort -R nums.txt -o rand_nums.txt

$ cat rand_nums.txt

```

```
1000
3.14
42
```

You can use `-o` for in-place editing as well, but the documentation gives this warning:

However, it is often safer to output to an otherwise-unused file, as data may be lost if the system crashes or `sort` encounters an I/O or other serious error while a file is being sorted in place. Also, `sort` with `--merge` (`-m`) can open the output file before reading all input, so a command like `cat F | sort -m -o F - G` is not safe as `sort` might start writing `F` before `cat` is done reading it.

Merge sort

The `-m` option is useful if you have one or more sorted input files and need a single sorted output file. Typically the use case is that you want to add newly obtained data to existing sorted data. In such cases, you can sort only the new data separately and then combine all the sorted inputs using the `-m` option. Here's a sample timing comparison between different combinations of sorted/unordered inputs.

```
$ shuf -n1000000 -i1-999999999999 > n1.txt
$ shuf -n1000000 -i1-999999999999 > n2.txt
$ sort -n n1.txt > n1_sorted.txt
$ sort -n n2.txt > n2_sorted.txt

$ time sort -n n1.txt n2.txt > op1.txt
real    0m1.010s
$ time sort -mn n1_sorted.txt <(sort -n n2.txt) > op2.txt
real    0m0.535s
$ time sort -mn n1_sorted.txt n2_sorted.txt > op3.txt
real    0m0.218s

$ diff -sq op1.txt op2.txt
Files op1.txt and op2.txt are identical
$ diff -sq op1.txt op3.txt
Files op1.txt and op3.txt are identical
```



You might wonder if you can improve the performance of a single large file using the `-m` option. By default, `sort` already uses the number of available processors to split the input and merge. You can use the `--parallel` option to customize this behavior.

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `sort` will ensure to add a final NUL character even if not present in the input.

```
$ printf 'cherry\0apple\0banana' | sort -z | cat -v
apple^@banana^@cherry^@
```

Further Reading

A few options like `--compress-program` and `--files0-from` aren't covered in this book. See [sort manual](#) for details and examples. See also:

- [unix.stackexchange](#): Scalability of sort for gigantic files
- [stackoverflow](#): Sort by last field when the number of fields varies
- [Arch wiki](#): locale
- [ShellHacks](#): locale and language settings

uniq

The `uniq` command identifies similar lines that are adjacent to each other. There are various options to help you filter unique or duplicate lines, count them, group them, etc.

Retain single copy of duplicates

This is the default behavior of the `uniq` command. If adjacent lines are the same, only the first copy will be displayed in the output.

```
# only the adjacent lines are compared to determine duplicates
# which is why you get 'red' twice in the output for this input
$ printf 'red\nred\nred\ngreen\nred\nblue\nblue' | uniq
red
green
red
blue
```

You'll need sorted input to make sure all the input lines are considered to determine duplicates. For some cases, `sort -u` is enough, like the example shown below:

```
# same as sort -u for this case
$ printf 'red\nred\nred\ngreen\nred\nblue\nblue' | sort | uniq
blue
green
red
```

Sometimes though, you may need to sort based on some specific criteria and then identify duplicates based on the entire line contents. Here's an example:

```
# can't use sort -n -u here
$ printf '2 balls\n13 pens\n2 pins\n13 pens\n' | sort -n | uniq
2 balls
2 pins
13 pens
```



`sort+uniq` won't be suitable if you need to preserve the input order as well. You can use alternatives like `awk`, `perl` and `huniq` for such cases.

```
# retain single copy of duplicates, maintain input order
$ printf 'red\nred\nred\ngreen\nred\nblue\nblue' | awk '!seen[$0]++'
red
green
blue
```

Duplicates only

The `-d` option will display only the duplicate entries. That is, only if a line is seen more than once.

```
$ cat purchases.txt
coffee
```

```
tea
washing powder
coffee
toothpaste
tea
soap
tea

$ sort purchases.txt | uniq -d
coffee
tea
```

To display all the copies of duplicates, use the `-D` option.

```
$ sort purchases.txt | uniq -D
coffee
coffee
tea
tea
tea
```

Unique only

The `-u` option will display only the unique entries. That is, only if a line doesn't occur more than once.

```
$ sort purchases.txt | uniq -u
soap
toothpaste
washing powder

# just a reminder that uniq works based on adjacent lines only
$ printf 'red\nred\nred\ngreen\nred\nblue\nblue' | uniq -u
green
red
```

Grouping similar lines

The `--group` options allows you to visually separate groups of similar lines with an empty line. This option can accept four values — `separate`, `prepend`, `append` and `both`. The default is `separate`, which adds a newline character between the groups. `prepend` will add a newline before the first group as well and `append` will add a newline after the last group. `both` combines `prepend` and `append` behavior.

```
$ sort purchases.txt | uniq --group
coffee
coffee

soap

tea
```

```
tea
tea

toothpaste

washing powder
```

The `--group` option cannot be used with `-c`, `-d`, `-D` or `-u` options. The `--all-repeated` alias for the `-D` option uses `none` as the default grouping. You can change that to `separate` or `prepend` values.

```
$ sort purchases.txt | uniq --all-repeated=prepend
```

```
coffee
coffee

tea
tea
tea
```

Prefix count

If you want to know how many times a line has been repeated, use the `-c` option. This will be added as a prefix.

```
$ sort purchases.txt | uniq -c
  2 coffee
  1 soap
  3 tea
  1 toothpaste
  1 washing powder
```

```
$ sort purchases.txt | uniq -dc
  2 coffee
  3 tea
```

The output of this option is usually piped to `sort` for ordering the output by the count value.

```
$ sort purchases.txt | uniq -c | sort -n
  1 soap
  1 toothpaste
  1 washing powder
  2 coffee
  3 tea
```

```
$ sort purchases.txt | uniq -c | sort -nr
  3 tea
  2 coffee
  1 washing powder
  1 toothpaste
  1 soap
```

Ignoring case

Use the `-i` option to ignore case while determining duplicates.

```
# depending on your locale, sort and sort -f can give the same results
$ printf 'cat\nbat\nCAT\ncar\nbat\nmat\nmoat' | sort -f | uniq -iD
bat
bat
cat
CAT
```

Partial match

`uniq` has three options to change the matching criteria to partial parts of the input line. These aren't as powerful as the `sort -k` option, but they do come in handy for some use cases.

The `-f` option allows you to skip first `N` fields. Field separation is based on one or more space/tab characters only. Note that these separators will still be part of the field contents, so this will not work with variable number of blanks.

```
# skip first field, works as expected since no. of blanks is consistent
$ printf '2 cars\n5 cars\n10 jeeps\n5 jeeps\n3 trucks\n' | uniq -f1 --group
2 cars
5 cars

10 jeeps
5 jeeps

3 trucks

# example with variable number of blanks
# 'cars' entries were identified as duplicates, but not 'jeeps'
$ printf '2 cars\n5 cars\n1 jeeps\n5 jeeps\n3 trucks\n' | uniq -f1
2 cars
1 jeeps
5 jeeps
3 trucks
```

The `-s` option allows you to skip first `N` characters (calculated as bytes).

```
# skip first character
$ printf '* red\n* green\n- green\n* blue\n= blue' | uniq -s1
* red
* green
* blue
```

The `-w` option restricts the comparison to the first `N` characters (calculated as bytes).

```
# compare only first 2 characters
$ printf '1) apple\n1) almond\n2) banana\n3) cherry' | uniq -w2
1) apple
2) banana
3) cherry
```


When these options are used simultaneously, the priority is `-f` first, then `-s` and finally `-w` option. Remember that blanks are part of the field content.

```
# skip first field
# then skip first two characters (including the blank character)
# use next two characters for comparison ('bl' and 'ch' in this example)
$ printf '2 @blue\n10 :black\n5 :cherry\n3 @chalk' | uniq -f1 -s2 -w2
2 @blue
5 :cherry
```



If a line doesn't have enough fields or characters to satisfy the `-f` and `-s` options respectively, a null string is used for comparison.

Specifying output file

`uniq` can accept filename as the source of input contents, but only a maximum of one file. If you specify another file, it will be used as the output file.

```
$ printf 'apple\napple\nbanana\ncherry\ncherry\ncherry' > ip.txt
$ uniq ip.txt op.txt

$ cat op.txt
apple
banana
cherry
```

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `uniq` will ensure to add a final NUL character even if not present in the input.

```
$ printf 'cherry\0cherry\0cherry\0apple\0banana' | uniq -z | cat -v
cherry^@apple^@banana^@
```



If grouping is specified, NUL will be used as the separator instead of the newline character.

Alternatives

Here's some alternate commands you can explore if `uniq` isn't enough to solve your task.

- [Dealing with duplicates](#) chapter from my **GNU awk** ebook
- [Dealing with duplicates](#) chapter from my **Perl one-liners** ebook
- `huniq` — remove duplicates from entire input contents, input order is maintained, supports count option as well

comm

The `comm` command finds common and unique lines between two sorted files. These results are formatted as a table with three columns and one or more of these columns can be suppressed as required.

Three column output

Consider the sample input files as shown below:

```
# side by side view of the sample files
# note that these files are already sorted
$ paste colors_1.txt colors_2.txt
Blue    Black
Brown   Blue
Orange  Green
Purple  Orange
Red     Pink
Teal    Red
White   White
```

By default, `comm` gives a tabular output with three columns:

- first column has lines unique to the first file
- second column has lines unique to the second file
- third column has lines common to both the files

The columns are separated by a tab character. Here's the output for the above sample files:

```
$ comm colors_1.txt colors_2.txt
      Black
      Blue
Brown
      Green
      Orange
      Pink
Purple
      Red
Teal
      White
```

You can change the column separator to a string of your choice using the `--output-delimiter` option. Here's an example:

```
# note that the input files need not have the same number of lines
$ comm <(seq 3) <(seq 2 5)
1
      2
      3
      4
      5

$ comm --output-delimiter=, <(seq 3) <(seq 2 5)
1
```

```
,,2
,,3
,4
,5
```



Collating order for `comm` should be same as the one used to `sort` the input files.



`--nocheck-order` option can be used for unsorted inputs. However, as per the documentation, this option "is not guaranteed to produce any particular output."

Suppressing columns

You can use one or more of the following options to suppress columns:

- `-1` to suppress lines unique to the first file
- `-2` to suppress lines unique to the second file
- `-3` to suppress lines common to both the files

Here's how the output looks like when you suppress one of the columns:

```
# suppress lines common to both the files
$ comm -3 colors_1.txt colors_2.txt
    Black
Brown
    Green
    Pink
Purple
Teal
```

Combining two of these options gives three useful solutions. `-12` will give you only the common lines.

```
$ comm -12 colors_1.txt colors_2.txt
Blue
Orange
Red
White
```

`-23` will give you the lines unique to the first file.

```
$ comm -23 colors_1.txt colors_2.txt
Brown
Purple
Teal
```

`-13` will give you the lines unique to the second file.

```
$ comm -13 colors_1.txt colors_2.txt
Black
Green
Pink
```

You can combine all the three options as well. Useful with the `--total` option to get only the count of lines for each of the three columns.

```
$ comm --total -123 colors_1.txt colors_2.txt
3      3      4      total
```

Duplicate lines

The number of duplicate lines in the common column will be minimum of the duplicate occurrences between the two files. Rest of the duplicate lines, if any, will be considered as unique to the file having the excess lines. Here's an example:

```
$ paste list_1.txt list_2.txt
apple  cherry
banana cherry
cherry mango
cherry papaya
cherry
cherry

# 'cherry' occurs only twice in the second file
# rest of the 'cherry' lines will be unique to the first file
$ comm list_1.txt list_2.txt
apple
banana
          cherry
          cherry
cherry
cherry
          mango
          papaya
```

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `comm` will ensure to add a final NUL character even if not present in the input.

```
$ comm -z -12 <(printf 'a\0b\0c') <(printf 'a\0c\0x') | cat -v
a^@c^@
```

Alternatives

Here's some alternate commands you can explore if `comm` isn't enough to solve your task. These alternatives do not require the input files to be sorted.

- [zet](#) — set operations on one or more input files
- [Comparing lines between files](#) section from my **GNU grep** ebook
- [Two file processing](#) chapter from my **GNU awk** ebook, both line and field based comparison
- [Two file processing](#) chapter from my **Perl one-liners** ebook, both line and field based comparison

join

The `join` command helps you to combine lines from two files based on a common field. This works best when the input is already sorted by that field.

Default join

By default, `join` combines two files based on the first field content (also referred as **key**). Only the lines with common keys will be part of the output.

The key field will be displayed first in the output (this distinction will come into play if the first field isn't the key). Rest of the line will have the remaining fields from the first and second files, in that order. One or more blanks (space or tab) will be considered as the input field separator and a single space will be used as the output field separator. If present, blank characters at the start of the input lines will be ignored.

```
# sample sorted input files
$ cat shopping_jan.txt
apple  10
banana 20
soap   3
tshirt 3
$ cat shopping_feb.txt
banana 15
fig     100
pen     2
soap    1

# combine common lines based on the first field
$ join shopping_jan.txt shopping_feb.txt
banana 20 15
soap 3 1
```

If a field value is present multiple times in the same input file, all possible combinations will be present in the output. As shown below, `join` will also ensure to add a final newline character even if not present in the input.

```
$ join <(printf 'a f1_x\na f1_y') <(printf 'a f2_x\na f2_y')
a f1_x f2_x
a f1_x f2_y
a f1_y f2_x
a f1_y f2_y
```



Note that the collating order used for `join` should be same as the one used to `sort` the input files. Use `join -i` to ignore case, similar to `sort -f` usage.



If the input files are not sorted, `join` will produce an error if there are unpairable lines. You can use the `--nocheck-order` option to ignore this error. However, as per the documentation, this option "is not guaranteed to produce any particular output."

Non-matching lines

By default, only the lines having common keys are part of the output. You can use the `-a` option to also include the non-matching lines from the input files. Use `1` and `2` as the argument for the first and second file respectively. You'll later see how to fill missing fields with a custom string.

```
# includes non-matching lines from the first file
$ join -a1 shopping_jan.txt shopping_feb.txt
apple 10
banana 20 15
soap 3 1
tshirt 3

# includes non-matching lines from both the files
$ join -a1 -a2 shopping_jan.txt shopping_feb.txt
apple 10
banana 20 15
fig 100
pen 2
soap 3 1
tshirt 3
```

If you use `-v` instead of `-a`, the output will have only the non-matching lines.

```
$ join -v1 shopping_jan.txt shopping_feb.txt
apple 10
tshirt 3

$ join -v1 -v2 shopping_jan.txt shopping_feb.txt
apple 10
fig 100
pen 2
tshirt 3
```

Change field separator

You can use the `-t` option to specify a single byte character as the field separator. The output field separator will be same as the value used for the `-t` option. Use `\0` to specify NUL as the separator. Empty string will cause entire input line content to be considered as keys. Depending on your shell you can use [ANSI-C quoting](#) to use escapes like `\t` instead of a literal tab character.

```
$ cat marks.csv
ECE,Raj,53
ECE,Joel,72
EEE,Moi,68
CSE,Surya,81
EEE,Raj,88
CSE,Moi,62
EEE,Tia,72
ECE,Om,92
CSE,Amy,67
```

```
$ cat dept.txt
CSE
ECE

# get all lines from marks.csv based on the first field keys in dept.txt
$ join -t, <(sort marks.csv) dept.txt
CSE,Amy,67
CSE,Moi,62
CSE,Surya,81
ECE,Joel,72
ECE,Om,92
ECE,Raj,53
```

Files with headers

Use the `--header` option to ignore first lines of both the input files from sorting consideration. Without this option, the `join` command might still work correctly if unpairable lines aren't found, but it is preferable to use `--header` when applicable. This option will also help when `--check-order` option is active.

```
$ cat report_1.csv
Name,Maths,Physics
Amy,78,95
Moi,88,75
Raj,67,76
$ cat report_2.csv
Name,Chemistry
Amy,85
Joel,78
Raj,72

$ join --check-order -t, report_1.csv report_2.csv
join: report_1.csv:2: is not sorted: Amy,78,95
$ join --check-order --header -t, report_1.csv report_2.csv
Name,Maths,Physics,Chemistry
Amy,78,95,85
Raj,67,76,72
```

Change key field

By default, the first field of both the input files are used to combine the lines. You can use `-1` and `-2` options followed by a field number to specify a different field number. You can use the `-j` option if the field number is the same for both the files.

Recall that the key field is the first field in the output. You'll later see how to customize the output field order.

```
$ cat names.txt
Amy
Raj
Tia
```

```
# combine based on second field of the first file
# and first field of the second file (default)
$ join -t, -1 2 <(sort -t, -k2,2 marks.csv) names.txt
Amy,CSE,67
Raj,ECE,53
Raj,EEE,88
Tia,EEE,72
```

Customize output field list

You can use the `-o` option to customize the fields required in the output and their order. Especially useful when the first field isn't the key. Each output field is specified as file number followed by a `.` character and then the field number. You can specify multiple fields separated by a `,` character. As a special case, you can use `0` to indicate the key field.

```
# output field order is 1st, 2nd and 3rd fields from the first file
$ join -t, -1 2 -o 1.1,1.2,1.3 <(sort -t, -k2,2 marks.csv) names.txt
CSE,Amy,67
ECE,Raj,53
EEE,Raj,88
EEE,Tia,72

# 1st field from the first file, 2nd field from the second file
# and then 2nd and 3rd fields from the first file
$ join --header -t, -o 1.1,2.2,1.2,1.3 report_1.csv report_2.csv
Name,Chemistry,Maths,Physics
Amy,85,78,95
Raj,72,67,76
```

Same number of output fields

If you use `auto` as the argument for the `-o` option, first line of both the input files will be used to determine the number of output fields. If the other lines have extra fields, they will be discarded.

```
$ join <(printf 'a 1 2\nb p q r') <(printf 'a 3 4\nb x y z')
a 1 2 3 4
b p q r x y z

$ join -o auto <(printf 'a 1 2\nb p q r') <(printf 'a 3 4\nb x y z')
a 1 2 3 4
b p q x y
```

If the other lines have lesser number of fields, the `-e` option will determine the string to be used as a filler (default is empty string).

```
$ join -o auto <(printf 'a 1 2\nb p') <(printf 'a 3 4\nb x')
a 1 2 3 4
b p x

$ join -o auto -e '-' <(printf 'a 1 2\nb p') <(printf 'a 3 4\nb x')
```



```
a 1 2 3 4
b p - x -
```

As promised earlier, here's an example of filling fields for non-matching lines:

```
$ join -o auto -a1 -e 'NA' shopping_jan.txt shopping_feb.txt
apple 10 NA
banana 20 15
soap 3 1
tshirt 3 NA

$ join -o auto -a1 -a2 -e 'NA' shopping_jan.txt shopping_feb.txt
apple 10 NA
banana 20 15
fig NA 100
pen NA 2
soap 3 1
tshirt 3 NA
```

Set operations

This section covers whole line set operations you can perform on already sorted input files. Equivalent `sort` and `uniq` solutions will also be mentioned as comments (useful for unsorted inputs). Assume that there are no duplicate lines within an input file.

These two sorted input files will be used for the examples to follow:

```
$ paste colors_1.txt colors_2.txt
Blue    Black
Brown   Blue
Orange  Green
Purple  Orange
Red     Pink
Teal    Red
White   White
```

Here's how you can get *union* and *symmetric difference* results. Recall that `-t ''` will cause entire input line content to be considered as keys.

```
# union
# unsorted input: sort -u colors_1.txt colors_2.txt
$ join -t '' -a1 -a2 colors_1.txt colors_2.txt
Black
Blue
Brown
Green
Orange
Pink
Purple
Red
Teal
White
```

```
# symmetric difference
# unsorted input: sort colors_1.txt colors_2.txt | uniq -u
$ join -t ' ' -v1 -v2 colors_1.txt colors_2.txt
Black
Brown
Green
Pink
Purple
Teal
```

Here's how you can get *intersection* and *difference* results. The equivalent `comm` solutions for sorted input is also mentioned in the comments.

```
# intersection, same as: comm -12 colors_1.txt colors_2.txt
# unsorted input: sort colors_1.txt colors_2.txt | uniq -d
$ join -t ' ' colors_1.txt colors_2.txt
Blue
Orange
Red
White

# difference, same as: comm -13 colors_1.txt colors_2.txt
# unsorted input: sort colors_1.txt colors_1.txt colors_2.txt | uniq -u
$ join -t ' ' -v2 colors_1.txt colors_2.txt
Black
Green
Pink

# difference, same as: comm -23 colors_1.txt colors_2.txt
# unsorted input: sort colors_1.txt colors_2.txt colors_2.txt | uniq -u
$ join -t ' ' -v1 colors_1.txt colors_2.txt
Brown
Purple
Teal
```

As mentioned before, `join` will display all the combinations if there are duplicate entries. Here's an example to show the differences between `sort`, `comm` and `join` solutions for displaying common lines:

```
$ paste list_1.txt list_2.txt
apple  cherry
banana cherry
cherry mango
cherry papaya
cherry
cherry

# only one entry per common line
$ sort list_1.txt list_2.txt | uniq -d
cherry
```

```
# minimum of 'no. of entries in file1' and 'no. of entries in file2'
$ comm -12 list_1.txt list_2.txt
cherry
cherry

# 'no. of entries in file1' multiplied by 'no. of entries in file2'
$ join -t ' ' list_1.txt list_2.txt
cherry
cherry
cherry
cherry
cherry
cherry
cherry
cherry
cherry
```

NUL separator

Use `-z` option if you want to use NUL character as the line separator. In this scenario, `join` will ensure to add a final NUL character even if not present in the input.

```
$ join -z <(printf 'a 1\0b x') <(printf 'a 2\0b y') | cat -v
a 1 2^@b x y^@
```

Alternatives

Here's some alternate commands you can explore if `join` isn't enough to solve your task. These alternatives do not require input to be sorted.

- [zet](#) — set operations on one or more input files
- [Comparing lines between files](#) section from my **GNU grep** ebook
- [Two file processing](#) chapter from my **GNU awk** ebook, both line and field based comparison
- [Two file processing](#) chapter from my **Perl one-liners** ebook, both line and field based comparison

nl

If the numbering options provided by `cat` isn't enough for you, `nl` might help you. Apart from options to customize the number formatting and the separator, you can also filter which lines should be numbered. Additionally, you can divide your input into sections and number them separately.

Default numbering

By default, `nl` will prefix line number and a tab character to every non-empty input lines. The default number formatting is 6 characters wide and right justified with spaces. Similar to `cat`, the `nl` command will concatenate multiple inputs.

```
# same as: cat -n greeting.txt fruits.txt nums.txt
$ nl greeting.txt fruits.txt nums.txt
 1 Hi there
 2 Have a nice day
 3 banana
 4 papaya
 5 mango
 6 3.14
 7 42
 8 1000

# example for input with empty lines, same as: cat -b
$ printf 'apple\n\nbanana\n\ncherry\n' | nl
 1 apple

 2 banana

 3 cherry
```

Number formatting

You can use the `-n` option to customize the number formatting. The available styles are:

- `rn` right justified with space fillers (default)
- `rz` right justified with leading zeros
- `ln` left justified with space fillers

```
# right justified with space fillers
$ nl -n'rn' greeting.txt
 1 Hi there
 2 Have a nice day

# right justified with leading zeros
$ nl -n'rz' greeting.txt
000001 Hi there
000002 Have a nice day

# left justified with space fillers
$ nl -n'ln' greeting.txt
```

```
1      Hi there
2      Have a nice day
```

Customize width

You can use the `-w` option to specify the width to be used for the numbers (default is `6`).

```
$ nl -w2 greeting.txt
1      Hi there
2      Have a nice day
```

Customize separator

By default, a tab character is used to separate the line number and the line content. You can use the `-s` option to specify your own custom string separator.

```
$ nl -w2 -s' ' greeting.txt
1 Hi there
2 Have a nice day

$ nl -w1 -s' --> ' greeting.txt
1 --> Hi there
2 --> Have a nice day
```

Starting number and increment

The `-v` option allows you to specify a different starting integer. Negative integer is also allowed.

```
$ nl -v10 greeting.txt
10 Hi there
11 Have a nice day

$ nl -v-1 fruits.txt
-1 banana
0 papaya
1 mango
```

The `-i` option allows you to specify a positive integer as the step value (default is `1`).

```
$ nl -w2 -s') ' -i2 greeting.txt fruits.txt nums.txt
1) Hi there
3) Have a nice day
5) banana
7) papaya
9) mango
11) 3.14
13) 42
15) 1000
```

Section wise numbering

If you organize your input with lines conforming to specific patterns, you can control their numbering separately. `nl` recognizes three types of sections with the following default patterns:

- `\:\:\:` as header
- `\:\:` as body
- `\:` as footer

These special lines will be replaced with an empty line after numbering. The numbering will be reset at the start of every section. Here's an example with multiple body sections:

```
$ cat body.txt
\:\:
Hi there
How are you
\:\:
banana
papaya
mango

$ nl -w1 -s' ' body.txt

1 Hi there
2 How are you

1 banana
2 papaya
3 mango
```

Here's an example with both header and body sections. By default, header and footer section lines are not numbered (you'll see options to enable them later).

```
$ cat header_body.txt
\:\:\:
Header
red
\:\:
Hi there
How are you
\:\:
banana
papaya
mango
\:\:\:
Header
green

$ nl -w1 -s' ' header_body.txt

Header
red
```

```
1 Hi there
2 How are you
```

```
1 banana
2 papaya
3 mango
```

```
Header
green
```

And here's an example with all the three types of sections:

```
$ cat all_sections.txt
```

```
\:\:\:
```

```
Header
```

```
red
```

```
\:\:
```

```
Hi there
```

```
How are you
```

```
\:\:
```

```
banana
```

```
papaya
```

```
mango
```

```
\:
```

```
Footer
```

```
$ nl -w1 -s' ' all_sections.txt
```

```
Header
```

```
red
```

```
1 Hi there
```

```
2 How are you
```

```
1 banana
```

```
2 papaya
```

```
3 mango
```

```
Footer
```

The `-b`, `-h` and `-f` options control which lines should be numbered for the three types of sections. Use `a` to number all lines of a particular section (other features will be discussed later).

```
$ nl -w1 -s' ' -ha -fa all_sections.txt
```

```
1 Header
```

```
2 red
```

```
1 Hi there
```

```
2 How are you
```

```
1 banana
2 papaya
3 mango

1 Footer
```

If you use the `-p` option, the numbering will not be reset on encountering a new section.

```
$ nl -p -w1 -s' ' all_sections.txt
```

```
Header
red

1 Hi there
2 How are you
```

```
3 banana
4 papaya
5 mango
```

```
Footer
```

```
$ nl -p -w1 -s' ' -ha -fa all_sections.txt
```

```
1 Header
2 red

3 Hi there
4 How are you
```

```
5 banana
6 papaya
7 mango
```

```
8 Footer
```

The `-d` option allows you to customize the two character pattern used for sections.

```
# pattern changed from \: to %=
```

```
$ cat body_sep.txt
```

```
%=%=
```

```
apple
```

```
banana
```

```
%=%=
```

```
red
```

```
green
```

```
$ nl -w1 -s' ' -d'%=' body_sep.txt
```

```
1 apple
2 banana
```



```
1 red
2 green
```

Section numbering criteria

As mentioned earlier, the `-b` , `-h` and `-f` options control which lines should be numbered for the three types of sections. These options accept the following arguments:

- `a` number all lines, including empty lines
- `t` number lines except empty ones (default for body sections)
- `n` do not number lines (default for header and footer sections)
- `pBRE` use basic regular expressions (BRE) to filter lines for numbering

If the input doesn't have special patterns to identify the different sections, it will be treated as if it has a single body section. Here's an example to include empty lines for numbering:

```
$ printf 'apple\n\nbanana\n\ncherry\n' | nl -w1 -s' ' -ba
1 apple
2
3 banana
4
5 cherry
```

The `-l` option controls how many consecutive empty lines should be considered as a single entry. Only the last empty line of such groupings will be numbered.

```
# only 2nd consecutive empty line will be considered for numbering
$ printf 'a\n\n\n\nb\n\nc' | nl -w1 -s' ' -ba -l2
1 a
2
3
4 b
5 c
```

Here's an example which uses regexp to identify lines to be numbered:

```
# number lines starting with 'c' or 't'
$ nl -w1 -s' ' -bp'^[ct]' purchases.txt
1 coffee
2 tea
   washing powder
3 coffee
4 toothpaste
5 tea
   soap
6 tea
```



See [Regular Expressions](#) chapter from my **GNU grep** ebook if you want to learn about regexp syntax and features.

WC

The `wc` command is useful to count the number of lines, words and characters for the given input(s).

Line, word and byte counts

By default, the `wc` command reports the number of lines, words and bytes (in that order). The byte count includes the newline characters, so you can use that as a measure of file size as well. Here's an example:

```
$ cat greeting.txt
Hi there
Have a nice day

$ wc greeting.txt
 2  6 25 greeting.txt
```

Wondering why there are leading spaces in the output? They help in aligning results for multiple files (discussed later).

Individual counts

Instead of the three default values, you can use options to get only the particular count(s) you are interested in. These options are:

- `-l` for line count
- `-w` for word count
- `-c` for byte count

```
$ wc -l greeting.txt
2 greeting.txt

$ wc -w greeting.txt
6 greeting.txt

$ wc -c greeting.txt
25 greeting.txt

$ wc -wc greeting.txt
 6 25 greeting.txt
```

With `stdin` data, you'll get only the count value (unless you use `-` for `stdin`). Useful for assigning the output to shell variables.

```
$ printf 'hello' | wc -c
5
$ printf 'hello' | wc -c -
5 -

$ lines=$(wc -l <greeting.txt)
$ echo "$lines"
2
```

Multiple files

If you pass multiple files to the `wc` command, the count values will be displayed separately for each file. You'll also get a summary at the end, which sums the respective count of all the input files.

```
$ wc greeting.txt nums.txt purchases.txt
 2  6 25 greeting.txt
 3  3 13 nums.txt
 8  9 57 purchases.txt
13 18 95 total
$ wc greeting.txt nums.txt purchases.txt | tail -n1
13 18 95 total

$ wc *[ck]*.csv
 9  9 101 marks.csv
 4  4  70 scores.csv
13 13 171 total
```

If you have NUL separated filenames (for example, output from `find -print0` , `grep -lZ` , etc), you can use the `--files0-from` option. This option accepts a file containing the NUL separated data (use `-` for `stdin`).

```
$ printf 'greeting.txt\0nums.txt' | wc --files0-from=-
2 6 25 greeting.txt
3 3 13 nums.txt
5 9 38 total
```

Character count

Use `-m` option instead of `-c` if the input has multibyte characters.

```
# byte count
$ printf 'αλεπού' | wc -c
12

# character count
$ printf 'αλεπού' | wc -m
6
```



Note that the current locale will affect the behavior of `-m` option.

```
$ printf 'αλεπού' | LC_ALL=C wc -m
12
```

Longest line length

You can use the `-L` to report the length of the longest line in the input (excluding the newline character of a line).

```
$ echo 'apple' | wc -L
5
```

```
# last line not ending with newline won't be a problem
$ printf 'apple\nbanana' | wc -L
6

$ wc -L sample.txt
26 sample.txt
$ wc -L <sample.txt
26
```

If multiple files are passed, the last line summary will show the maximum length among the given inputs.

```
$ wc -L greeting.txt nums.txt purchases.txt
15 greeting.txt
 4 nums.txt
14 purchases.txt
15 total
```

Corner cases

Line count is based on the number of newline characters. So, if the last line of the input doesn't end with the newline character, it won't be counted.

```
$ printf 'good\nmorning\n' | wc -l
2

$ printf 'good\nmorning' | wc -l
1

$ printf '\n\n\n' | wc -l
3
```

Word count is based on whitespace separation. You'll have to pre-process the input if you do not want certain non-whitespace characters to influence the results.

```
$ echo 'apple ; banana ; cherry' | wc -w
5

# remove characters other than alphabets and whitespace
$ echo 'apple ; banana ; cherry' | tr -cd 'a-zA-Z[:space:]'
apple banana cherry
$ echo 'apple ; banana ; cherry' | tr -cd 'a-zA-Z[:space:]' | wc -w
3

# allow numbers as well
$ echo '2 : apples ;' | tr -cd '[:alnum:][:space:]' | wc -w
2
```

`-L` won't count non-printable characters and tabs are converted to equivalent spaces. Multibyte characters will each be counted as `1` (depending on the locale, they might become non-printable too).

```

# tab characters can occupy up to 8 columns
$ printf '\t' | wc -L
8
$ printf 'a\tb' | wc -L
9

# example for non-printable character
$ printf 'a\34b' | wc -L
2

# multibyte characters are counted as 1 each in supported locales
$ printf 'αλεπού' | wc -L
6
# non-supported locales can cause them to be treated as non-printable
$ printf 'αλεπού' | LC_ALL=C wc -L
0

```

`-m` and `-L` options count [grapheme clusters](#) differently.

```

$ printf 'caĝe' | wc -m
5

$ printf 'caĝe' | wc -L
4

```

split

The `split` command is useful to divide the input into smaller parts based on number of lines, bytes, file size, etc. You can also execute another command on the divided parts before saving the results. An example use case is sending a large file as multiple parts as a workaround for online transfer size limits.



Since a lot of output files will be generated in this chapter (often with same filenames), remove these files after every illustration.

Default split

By default, the `split` command divides the input `1000` lines at a time. Newline character is the default line separator. You can pass a single file or `stdin` data as the input. Use `cat` if you need to concatenate multiple input sources.

By default, the output files will be named `xaa` , `xab` , `xac` and so on (where `x` is the prefix). If the filenames are exhausted, two more letters will be appended and the pattern will continue as needed. If the number of input lines is not evenly divisible, the last file will contain less than `1000` lines.

```
# divide input 1000 lines at a time
$ seq 10000 | split

# output filenames
$ ls x*
xaa xab xac xad xae xaf xag xah xai xaj

# preview of some of the output files
$ head -n1 xaa xab xae xaj
==> xaa <==
1

==> xab <==
1001

==> xae <==
4001

==> xaj <==
9001

$ rm x*
```



As mentioned earlier, remove the output files after every illustration.

Change number of lines

You can use the `-l` option to change the number of lines to be saved in each output file.

```
# maximum of 3 lines at a time
$ split -l3 purchases.txt

$ head x*
==> xaa <==
coffee
tea
washing powder

==> xab <==
coffee
toothpaste
tea

==> xac <==
soap
tea
```

Split by byte count

The `-b` option allows you to split the input by number of bytes. Similar to line based splitting, you can always reconstruct the input by concatenating the output files. This option also accepts suffixes such as `K` for 1024 bytes, `KB` for 1000 bytes, `M` for 1024 * 1024 bytes and so on.

```
# maximum of 15 bytes at a time
$ split -b15 greeting.txt

$ head x*
==> xaa <==
Hi there
Have a
==> xab <==
  nice day

# when you concatenate the output files, you'll the original input
$ cat x*
Hi there
Have a nice day
```

The `-C` option is similar to the `-b` option, but it will try to break on line boundaries if possible. The break will happen before the given byte limit. Here's an example where input lines do not exceed the given byte limit:

```
$ split -C20 purchases.txt

$ head x*
==> xaa <==
coffee
tea
```

```

==> xab <==
washing powder

==> xac <==
coffee
toothpaste

==> xad <==
tea
soap
tea

$ wc -c x*
11 xaa
15 xab
18 xac
13 xad
57 total

```

If a line exceeds the given limit, it will be broken down into multiple parts:

```

$ printf 'apple\nbanana\n' | split -C4

$ head x*
==> xaa <==
appl
==> xab <==
e

==> xac <==
bana
==> xad <==
na

$ cat x*
apple
banana

```

Divide based on file size

The `-n` option has several features. If you pass only a numeric argument `N`, the given input file will be divided into `N` chunks. The output files will be roughly the same size.

```

# divide the file into 2 parts
$ split -n2 purchases.txt
$ head x*
==> xaa <==
coffee
tea
washing powder
co

```



```

==> xab <==
ffee
toothpaste
tea
soap
tea

# the two output files are roughly the same size
$ wc x*
 3  5 28 xaa
 5  5 29 xab
 8 10 57 total

```



Since the division is based on file size, `stdin` data cannot be used.

```

$ seq 6 | split -n2
split: -: cannot determine file size

```

By using `K/N` as the argument, you can view the `K` th chunk of `N` parts on `stdout` . No output file will be created in this scenario.

```

# divide the input into 2 parts
# view only the 1st chunk on stdout
$ split -n1/2 greeting.txt
Hi there
Hav

```

To avoid splitting a line, use `l/` as a prefix. Quoting from the [manual](#):

For `l` mode, chunks are approximately `input size / N` . The input is partitioned into `N` equal sized portions, with the last assigned any excess. If a line starts within a partition it is written completely to the corresponding file. Since lines or records are not split even if they overlap a partition, the files written can be larger or smaller than the partition size, and even empty if a line/record is so long as to completely overlap the partition.

```

# divide input into 2 parts, don't split lines
$ split -nl/2 purchases.txt
$ head x*
==> xaa <==
coffee
tea
washing powder
coffee

==> xab <==
toothpaste
tea
soap
tea

```

Here's an example to view `K` th chunk without splitting lines:

```
# 2nd chunk of 3 parts, don't split lines
$ split -nl/2/3 sample.txt
7) Believe it
8)
9) banana
10) papaya
11) mango
```

Interleaved lines

The `-n` option will also help you create output files with interleaved lines. Since this is based on the line separator and not file size, `stdin` data can also be used. Use `r/` prefix to enable this feature.

```
# two parts, lines distributed in round robin fashion
$ seq 5 | split -nr/2

$ head x*
==> xaa <==
1
3
5

==> xab <==
2
4
```

Here's an example to view `K` th chunk:

```
$ split -nr/1/3 sample.txt
1) Hello World
4) How are you
7) Believe it
10) papaya
13) Much ado about nothing
```

Custom line separator

You can use the `-t` option to specify a single byte character as the line separator. Use `\0` to specify NUL as the separator. Depending on your shell you can use [ANSI-C quoting](#) to use escapes like `\t` instead of a literal tab character.

```
$ printf 'apple\nbanana\n;mango\npapaya\n' | split -t';' -l1

$ head x*
==> xaa <==
apple
banana
;

==> xab <==
```

```
mango
papaya
```

Customize filenames

As seen earlier, `x` is the default prefix for output filenames. To change this prefix, pass an argument after the input source.

```
# choose prefix as 'op_' instead of 'x'
$ split -l1 greeting.txt op_

$ head op_*
==> op_aa <==
Hi there

==> op_ab <==
Have a nice day
```

The `-a` option controls the length of the suffix. You'll get an error if this length isn't enough to cover all the output files. In such a case, you'll still get output files that can fit within the given length.

```
$ seq 10 | split -l1 -a1
$ ls x*
xa xb xc xd xe xf xg xh xi xj
$ rm x*

$ seq 10 | split -l1 -a3
$ ls x*
xaaa xaab xaac xaad xaae xaaf xaag xaah xaai xaaj
$ rm x*

$ seq 100 | split -l1 -a1
split: output file suffixes exhausted
$ ls x*
xa xc xe xg xi xk xm xo xq xs xu xw xy
xb xd xf xh xj xl xn xp xr xt xv xx xz
```

You can use the `-d` option to use numeric suffixes, starting from `00` (length can be changed using the `-a` option). You can use the long option `--numeric-suffixes` to specify a different starting number.

```
$ seq 10 | split -l1 -d
$ ls x*
x00 x01 x02 x03 x04 x05 x06 x07 x08 x09
$ rm x*

$ seq 10 | split -l2 --numeric-suffixes=10
$ ls x*
x10 x11 x12 x13 x14
```

Use `-x` and `--hex-suffixes` options for hexadecimal numbering.

```
$ seq 10 | split -l1 --hex-suffixes=8
$ ls x*
x08 x09 x0a x0b x0c x0d x0e x0f x10 x11
```

You can use the `--additional-suffix` option to add a constant string at the end of filenames.

```
$ seq 10 | split -l2 -a1 --additional-suffix='.log'
$ ls x*
xa.log xb.log xc.log xd.log xe.log
$ rm x*

$ seq 10 | split -l2 -a1 -d --additional-suffix='.txt' - num_
$ ls num_*
num_0.txt num_1.txt num_2.txt num_3.txt num_4.txt
```

Exclude empty files

You can sometimes end up with empty files. For example, trying to split into more parts than possible with the given criteria. In such cases, you can use the `-e` option to prevent empty files in the output. The `split` command will ensure that the filenames are sequential even if files in the middle are empty.

```
# 'xac' is empty in this example
$ split -nl/3 greeting.txt
$ head x*
==> xaa <==
Hi there

==> xab <==
Have a nice day

==> xac <==

$ rm x*

# prevent empty files
$ split -e -nl/3 greeting.txt
$ head x*
==> xaa <==
Hi there

==> xab <==
Have a nice day
```

Process parts through another command

The `--filter` option will allow you to apply another command on the intermediate `split` results before saving the output files. Use `$FILE` to refer to the output filename of the intermediate parts. Here's an example of compressing the results:

```
$ split -l1 --filter='gzip > $FILE.gz' greeting.txt

$ ls x*
xaa.gz  xab.gz

$ zcat xaa.gz
Hi there
$ zcat xab.gz
Have a nice day
```

Here's an example of ignoring the first line of the results:

```
$ cat body_sep.txt
%=%=
apple
banana
%=%=
red
green

$ split -l3 --filter='tail -n +2 > $FILE' body_sep.txt

$ head x*
==> xaa <==
apple
banana

==> xab <==
red
green
```

csplit

The `csplit` command is useful to divide the input into smaller parts based on line numbers and regular expression patterns. Similar to `split`, this command also supports customizing output filenames.



Since a lot of output files will be generated in this chapter (often with same filenames), remove these files after every illustration.

Split on Nth line

You can split the input into two based on a particular line number. To do so, specify the line number after the input source (filename or `stdin` data). The first output file will have the input lines *before* the given line number and the second output file will have the rest of the contents.

By default, the output files will be named `xx00`, `xx01`, `xx02` and so on (where `xx` is the prefix). The numerical suffix will automatically use more digits if needed. You'll see examples with more than two output files later.

```
# split input into two based on line number 4
$ seq 10 | csplit - 4
6
15

# first output file will have the first 3 lines
# second output file will have the rest
$ head xx*
==> xx00 <==
1
2
3

==> xx01 <==
4
5
6
7
8
9
10

$ rm xx*
```



As seen in the example above, `csplit` will also display the number of bytes written for each output file. You can use the `-q` option to suppress this message.



As mentioned earlier, remove the output files after every illustration.

Split on regexp

You can also split the input based on a line matching the given regular expression. The output produced will vary based on `//` or `%` delimiters being used to surround the regexp.

When `/regexp/` is used, output is similar to the line number based splitting. The first output file will have the input lines *before* the first occurrence of a line matching the given regexp and the second output file will have the rest of the contents.

```
# match a line containing 't' followed by zero or more characters and then 'p'
# 'toothpaste' is the only match for this input file
$ csplit -q purchases.txt '/t.*p/'

$ head xx*
==> xx00 <==
coffee
tea
washing powder
coffee

==> xx01 <==
toothpaste
tea
soap
tea
```

When `%regexp%` is used, the lines occurring before the matching line won't be part of the output. Only the line matching the given regexp and the rest of the contents will be part of the single output file.

```
$ csplit -q purchases.txt '%t.*p%'

$ cat xx00
toothpaste
tea
soap
tea
```



You'll get an error if the given regexp isn't found in the input.

```
$ csplit -q purchases.txt '/xyz/'
csplit: '/xyz/': match not found
```



See [Regular Expressions](#) chapter from my **GNU grep** ebook if you want to learn about regexp syntax and features.

Regexp offset

You can also provide offset numbers that'll affect where the matching line and its surrounding lines should be placed. When the offset is greater than zero, the split will happen that many lines after the matching line. The default offset is zero.

```
# when the offset is '1', matching line will be part of the first file
$ csplit -q purchases.txt '/t.*p/1'
$ head xx*
==> xx00 <==
coffee
tea
washing powder
coffee
toothpaste

==> xx01 <==
tea
soap
tea

$ rm xx*

# matching line and 1 line after won't be part of the output
$ csplit -q purchases.txt '%t.*p%2'
$ cat xx00
soap
tea
```

When the offset is less than zero, the split will happen that many lines before the matching line.

```
# 2 lines before the matching line will be part of the second file
$ csplit -q purchases.txt '/t.*p/-2'
$ head xx*
==> xx00 <==
coffee
tea

==> xx01 <==
washing powder
coffee
toothpaste
tea
soap
tea
```



You'll get an error if the offset goes beyond the number of lines available in the input.

```
$ csplit -q purchases.txt '/t.*p/5'
csplit: '/t.*p/5': line number out of range
```



```
$ csplit -q purchases.txt '/t.*p/-5'
csplit: '/t.*p/-5': line number out of range
```

Repeat split

You can perform line number and regexp based split more than once by adding `{N}` argument after the pattern. Default behavior examples seen so far is same as specifying `{0}`. Any number greater than zero will result in that many more splits.

```
# {1} means split one time more than the default split
# so, two splits in total and three output files
# in this example, split happens on 4th and 8th line numbers
$ seq 10 | csplit -q - 4 '{1}'

$ head xx*
==> xx00 <==
1
2
3

==> xx01 <==
4
5
6
7

==> xx02 <==
8
9
10
```

Here's an example with regexp:

```
$ cat log.txt
--> warning 1
a,b,c,d
42
--> warning 2
x,y,z
--> warning 3
4,3,1

# split on 3rd (2+1) occurrence of a line containing 'warning'
$ csplit -q log.txt '%warning%' '{2}'
$ cat xx00
--> warning 3
4,3,1
```

As a special case, you can use `{*}` to repeat the split until the input is exhausted. This is especially useful with the `/regexp/` form of splitting. Here's an example:

```
# split on all lines matching 'paste' or 'powder'
$ csplit -q purchases.txt '/paste\|powder/' '{*}'
$ head xx*
==> xx00 <==
coffee
tea

==> xx01 <==
washing powder
coffee

==> xx02 <==
toothpaste
tea
soap
tea
```



You'll get an error if the repeat count goes beyond the number of matches possible with the given input.

```
$ seq 10 | csplit -q - 4 '{2}'
csplit: '4': line number out of range on repetition 2

$ csplit -q purchases.txt '/tea/' '{4}'
csplit: '/tea/': match not found on repetition 3
```

Keep files on error

By default, `csplit` will remove the created output files if there's an error or a signal that causes the command to stop. You can use the `-k` option to keep such files. One use case is line number based splitting with the `{*}` modifier.

```
$ seq 10 | csplit -q - 4 '{*}'
csplit: '4': line number out of range on repetition 2
$ ls xx*
ls: cannot access 'xx*': No such file or directory

# -k option will allow you to retain the created files
$ seq 10 | csplit -qk - 4 '{*}'
csplit: '4': line number out of range on repetition 2
$ head xx*
==> xx00 <==
1
2
3

==> xx01 <==
4
5
```

```
6
7

==> xx02 <==
8
9
10
```

Suppress matched lines

The `--suppress-matched` option will suppress the lines matching the split condition.

```
$ seq 5 | csplit -q --suppress-matched - 3
# 3rd line won't be part of the output
$ head xx*
==> xx00 <==
1
2

==> xx01 <==
4
5

$ rm xx*

$ seq 10 | csplit -q --suppress-matched - 4 '{1}'
# 4th and 8th lines won't be part of the output
$ head xx*
==> xx00 <==
1
2
3

==> xx01 <==
5
6
7

==> xx02 <==
9
10
```

Here's an example with regexp based split:

```
$ csplit -q --suppress-matched purchases.txt '/soap|powder/' '{*}'
# lines matching 'soap' or 'powder' won't be part of the output
$ head xx*
==> xx00 <==
coffee
tea
```

```
==> xx01 <==
coffee
toothpaste
tea

==> xx02 <==
tea
```



Suppressing matched lines for regexp based split other than `{*}` usage doesn't give expected results. See [this bug report](#) for more details. This bug has been fixed in [coreutils version 9.0](#).

```
$ seq 11 14 | csplit -q --suppress-matched - '/3/'
# the matching line wasn't suppressed
$ head xx*
==> xx00 <==
11
12

==> xx01 <==
13
14

$ rm xx*

$ seq 11 16 | csplit -q --suppress-matched - '/[35]/' '{1}'
# the first matching line was correctly suppressed
# but the second matching line wasn't suppressed
$ head xx*
==> xx00 <==
11
12

==> xx01 <==
14

==> xx02 <==
15
16
```

Exclude empty files

There are various cases that can result in empty output files. For example, first or last line matching the given split condition. Another possibility is `--suppress-matched` option combined with consecutive lines matching during multiple splits. Here's an example:

```
$ csplit -q --suppress-matched purchases.txt '/coffee|tea/' '{*}'

$ head xx*
==> xx00 <==
```

```

==> xx01 <==

==> xx02 <==
washing powder

==> xx03 <==
toothpaste

==> xx04 <==
soap

==> xx05 <==

```

You can use the `-z` option to exclude empty files from the output. The suffix numbering will be automatically adjusted in such cases.

```

$ csplit -qz --suppress-matched purchases.txt '/coffee\|tea/' '{*}'

$ head xx*
==> xx00 <==
washing powder

==> xx01 <==
toothpaste

==> xx02 <==
soap

```

Customize filenames

As seen earlier, `xx` is the default prefix for output filenames. Use the `-f` option to change this prefix.

```

$ seq 4 | csplit -q -f'num_' - 3

$ head num_*
==> num_00 <==
1
2

==> num_01 <==
3
4

```

The `-n` option controls the length of the numeric suffix. The suffix length will automatically increment if filenames are exhausted.

```

$ seq 4 | csplit -q -n1 - 3
$ ls xx*
xx0  xx1
$ rm xx*

```

```
$ seq 4 | csplit -q -n3 - 3
$ ls xx*
xx000  xx001
```

The `-b` option allows you to control the suffix using `printf` formatting. Quoting from the [manual](#):

When this option is specified, the suffix string must include exactly one `printf(3)` -style conversion specification, possibly including format specification flags, a field width, a precision specifications, or all of these kinds of modifiers. The format letter must convert a binary unsigned integer argument to readable form. The format letters `d` and `i` are aliases for `u`, and the `u`, `o`, `x`, and `X` conversions are allowed.

Here's some examples:

```
# hexadecimal numbering
# minimum two digits, zero filled
$ seq 100 | csplit -q -b'%02x' - 3 '{20}'
$ ls xx*
xx00  xx02  xx04  xx06  xx08  xx0a  xx0c  xx0e  xx10  xx12  xx14
xx01  xx03  xx05  xx07  xx09  xx0b  xx0d  xx0f  xx11  xx13  xx15
$ rm xx*

# custom prefix and suffix around decimal numbering
# default minimum is single digit
$ seq 20 | csplit -q -f'num_' -b'%d.txt' - 3 '{4}'
$ ls num_*
num_0.txt  num_1.txt  num_2.txt  num_3.txt  num_4.txt  num_5.txt
```



Note that the `-b` option will override the `-n` option. See `man 3 printf` for more details about the formatting options.

expand and unexpand

These two commands will help you convert tabs to spaces and vice versa. Both these commands support options to customize the width of tab stops and which occurrences should be converted.

Default expand

The `expand` command converts tab characters to space characters. The default expansion aligns at multiples of `8` columns (calculated in terms of bytes).

```
# sample stdin data
$ printf 'apple\tbanana\tcherry\n\tb\tc\n' | cat -T
apple^Ibanana^Icherry
a^Ib^Ic
# 'apple' = 5 bytes, \t converts to 3 spaces
# 'banana' = 6 bytes, \t converts to 2 spaces
# 'a' and 'b' = 1 byte, \t converts to 7 spaces
$ printf 'apple\tbanana\tcherry\n\tb\tc\n' | expand
apple  banana  cherry
a      b      c

# 'αλε' = 6 bytes, \t converts to 2 spaces
$ printf 'αλε\tπού\n' | expand
αλε   πού
```

Here's an example with strings of size `7` and `8` bytes before the tab character:

```
$ printf 'deviate\treached\nbackdrop\toverhang\n' | expand
deviate reached
backdrop      overhang
```

The `expand` command also considers backspace characters to determine the number of spaces needed.

```
# sample input with a backspace character
$ printf 'cart\bd\tbard\n' | cat -t
cart^Hd^Ibard

# 'card' = 4 bytes, \t converts to 4 spaces
$ printf 'cart\bd\tbard\n' | expand
card    bard
$ printf 'cart\bd\tbard\n' | expand | cat -t
cart^Hd    bard
```



`expand` will concatenate multiple files passed as input source, so `cat` will not be needed for such cases.

Expand only initial tabs

You can use the `-i` option to convert only the tab characters present at the start of a line. The first occurrence of a character that is not a tab or space character will stop the expansion.

```
# 'a' present at the start of line is not a tab/space character
# so no tabs are expanded for this input
$ printf 'a\tb\tc\n' | expand -i | cat -T
a^Ib^Ic

# the first \t gets expanded here, 'a' stops further expansion
$ printf '\ta\tb\tc\n' | expand -i | cat -T
a^Ib^Ic

# first two \t gets expanded here, 'a' stops further expansion
# presence of space characters will not stop the expansion
$ printf '\t \ta\tb\tc\n' | expand -i | cat -T
a^Ib^Ic
```

Customize tab stop width

You can use the `-t` option to control the expansion width. Default is `8` as seen in the previous examples.

This option provides various features. Here's an example where all the tab characters are converted equally to the given width:

```
$ cat -T code.py
def compute(x, y):
^Iif x > y:
^I^Iprint('hello')
^Ielse:
^I^Iprint('bye')

$ expand -t 2 code.py
def compute(x, y):
    if x > y:
        print('hello')
    else:
        print('bye')
```

You can provide multiple widths separated by a comma character. In such a case, the given widths determine the stop locations for those many tab characters. These stop values refer to absolute positions from the start of the line, not the number of spaces they can expand to. Rest of the tab characters will be expanded to a single space character.

```
# first tab character can expand till 3rd column
# second tab character can expand till 7th column
# rest of the tab characters will be expanded to single space
$ printf 'a\tb\tc\td\te\n' | expand -t 3,7
a  b   c d e

# here's two more examples with the same specification as above
# second tab expands to two spaces to end at 7th column
$ printf 'a\tbb\tc\td\te\n' | expand -t 3,7
a  bb  c d e

# second tab expands to single space since it goes beyond 7th column
```



```
$ printf 'a\tbbbbbbb\tc\td\te\n' | expand -t 3,7
a  bbbbbbbb c d e
```

If you prefix a `/` character to the last width, the remaining tab characters will use multiple of this position instead of single space default.

```
# first tab character can expand till 3rd column
# remaining tab characters can expand till 7/14/21/etc
$ printf 'a\tb\tc\td\te\tf\tg\n' | expand -t 3,/7
a  b  c      d      e      f      g
```

```
# first tab character can expand till 3rd column
# second tab character can expand till 7th column
# remaining tab characters can expand till 10/15/20/etc
$ printf 'a\tb\tc\td\te\tf\tg\n' | expand -t 3,7,/5
a  b   c   d   e   f   g
```

If you use `+` instead of `/` as the prefix for the last width, the multiple calculation will use the second last width as an offset.

```
# first tab character can expand till 3rd column
# 3+7=10, so remaining tab characters can expand till 10/17/24/etc
$ printf 'a\tb\tc\td\te\tf\tg\n' | expand -t 3,+7
a  b      c      d      e      f      g

# first tab character can expand till 3rd column
# second tab character can expand till 7th column
# 7+5=12, so remaining tab characters can expand till 12/17/22/etc
$ printf 'a\tb\tc\td\te\tf\tg\n' | expand -t 3,7,+5
a  b  c      d      e      f      g
```

Default unexpand

By default, the `unexpand` command converts initial blank (space or tab) characters to tabs. The first occurrence of a non-blank character will stop the conversion. By default, every 8 columns worth of blanks is converted to a tab.

[illegible]

```
# input has 4 spaces and a tab character (that expands till 8th column)
# output will have a single tab character at the start
$ printf '    \ta b\n' | unexpand | cat -T
^Ia b
```



The current `locale` determines which characters are considered as blanks. Also, `unexpand` will concatenate multiple files passed as input source, so `cat` will not be needed for such cases.

Unexpand all blanks

The `-a` option will allow you to convert all sequences of two or more blanks at tab boundaries. Here's some examples:

```
# default unexpand stops at first non-blank character
$ printf '    a    b    c\n' | unexpand | cat -T
^Ia    b    c

# -a option will convert all sequences of blanks at tab boundaries
$ printf '    a    b    c\n' | unexpand -a | cat -T
^Ia^Ib^Ic

# only two or more consecutive blanks are considered for conversion
$ printf 'riddled reached\n' | unexpand -a | cat -T
riddled reached
$ printf 'riddle reached\n' | unexpand -a | cat -T
riddle^Ireached

# blanks at non-tab boundaries won't be converted
$ printf 'oh hi hello\n' | unexpand -a | cat -T
oh hi^Ihello
```

The `unexpand` command also considers backspace characters to determine the tab boundary.

```
# 'card' = 4 bytes, so the 4 spaces gets converted to a tab
$ printf 'card\bd    bard\n' | unexpand -a | cat -T
card^Ibard
$ printf 'card\bd    bard\n' | unexpand -a | cat -t
card^Hd^Ibard
```

Change tab stop width

The `-t` option has the same features as seen with `expand` command. The `-a` option is also implied when this option is used.

Here's an example of changing the tab stop width to `2` :

```
$ printf '\ta\n\t\tb\n' | expand -t 2
a
  b

$ printf '\ta\n\t\tb\n' | expand -t 2 | unexpand -t 2 | cat -T
```

```
^Ia
^I^Ib
```

Here's some examples for multiple tab widths:

```
$ printf 'a\tb\tc\td\te\n' | expand -t 3,7
a  b   c d e
$ printf 'a  b   c d e\n' | unexpand -t 3,7 | cat -T
a^Ib^Ic d e
$ printf 'a\tb\tc\td\te\n' | expand -t 3,7 | unexpand -t 3,7 | cat -T
a^Ib^Ic d e

$ printf 'a\tb\tc\td\te\tf\n' | expand -t 3,/7
a  b   c       d       e       f
$ printf 'a  b   c       d       e       f\n' | unexpand -t 3,/7 | cat -T
a^Ib^Ic^Id^Ie^If

$ printf 'a\tb\tc\td\te\tf\n' | expand -t 3,+7
a  b       c       d       e       f
$ printf 'a  b       c       d       e       f\n' | unexpand -t 3,+7 | cat -T
a^Ib^Ic^Id^Ie^If
```

basename and dirname

These handy commands allow you to extract filenames and directory portions of the given paths. You could also use [Parameter Expansion](#) or `cut`, `sed`, `awk`, etc for such purposes. The advantage is that these commands will also handle corner cases like trailing slashes and there are handy features like removing file extensions.

Extract filename from path

By default, the `basename` command will remove the leading directory component from the given path argument. Any trailing slashes will be removed before determining the portion to be extracted.

```
$ basename /home/learnbyexample/example_files/scores.csv
scores.csv

# quote the arguments when needed
$ basename 'path with spaces/report.log'
report.log

# one or more trailing slashes will not affect the output
$ basename /home/learnbyexample/example_files/
example_files
```

If there's no leading directory component or if slash alone is the input, the argument will be returned as is after removing any trailing slashes.

```
$ basename filename.txt
filename.txt

$ basename /
/
```

Remove file extension

You can use the `-s` option to remove a suffix from the filename. Usually used to remove the file extension.

```
$ basename -s'.csv' /home/learnbyexample/example_files/scores.csv
scores

$ basename -s'_2' final_report.txt_2
final_report.txt

$ basename -s'.tar.gz' /backups/jan_2021.tar.gz
jan_2021

$ basename -s'.txt' purchases.txt.txt
purchases.txt

# -s will be ignored if it would have resulted in empty output
$ basename -s'report' /backups/report
```

```
report
```

You can also pass the suffix to be removed after the path argument, but the `-s` option is preferred as it makes the intention clearer and works for multiple path arguments.

```
$ basename example_files/scores.csv .csv
scores
```

Remove filename from path

By default, the `dirname` command removes the trailing path component (after removing any trailing slashes).

```
$ dirname /home/learnbyexample/example_files/scores.csv
/home/learnbyexample/example_files

# one or more trailing slashes will not affect the output
$ dirname /home/learnbyexample/example_files/
/home/learnbyexample
```

Multiple arguments

The `dirname` command accepts multiple path arguments by default. The `basename` command requires `-a` or `-s` (which implies `-a`) to work with multiple arguments.

```
$ basename -a /backups/jan_2021.tar.gz /home/learnbyexample/report.log
jan_2021.tar.gz
report.log

# -a is implied when -s is used
$ basename -s'.txt' logs/purchases.txt logs/report.txt
purchases
report

$ dirname /home/learnbyexample/example_files/scores.csv ../report/backups/
/home/learnbyexample/example_files
../report
```

Combining basename and dirname

You can use shell features like command substitution to combine the effects of `basename` and `dirname` commands.

```
# extract the second last path component
$ basename $(dirname /home/learnbyexample/example_files/scores.csv)
example_files
```

NUL separator

Use `-z` option if you want to use NUL character as the output path separator.

```
$ basename -zs'.txt' logs/purchases.txt logs/report.txt | cat -v
purchases^@report^@

$ basename -z logs/purchases.txt | cat -v
purchases.txt^@

$ dirname -z example_files/scores.csv ../report/backups/ | cat -v
example_files^@../report^@
```

What next?

Hope you've found this book interesting and useful.

There are plenty of general purpose and specialized text processing tools. Here's a list of books I've written (or currently working upon):

- [GNU grep and ripgrep](#)
- [GNU sed](#)
- [GNU awk](#)
- [Perl one-liners](#)
- [Ruby one-liners](#)
- [Command line text processing with Rust tools](#)

See my [curated list on cli text processing](#) for even more tools and resources.